TACC

TEXAS ADVANCED COMPUTING CENTER

WWW.TACC.UTEXAS.EDU

# OpenMP Affinity Concepts

CSC Summer Institute

**October 9-11, 2023**

Kent Milfeld (TACC)

Emanuele Vitali (CSC)

Slides at tinyurl.com/ adv-omp-2023

# What you will learn

- The basics of Affinity →
  - affinity fundamentals, the affinity mask, and how it is used
- How to acquire affinity mask (and use *amask* utility)
- How to use OpenMP  Affinity:

  (proc_bind and places of OpenMP)

# Sections

- Motivation
- Affinity fundamentals, the affinity mask, and how it is used
- OpenMP  Affinity:  proc_bind and places
- Showing Mask with amask utility

# Motivation

- HPC Nodes: many cores, sockets, SMT* (Simultaneous Multi-Threading on a core), NUMA
- HPC User's approach to executing an application is too often:
  Get app to compile (or use site-installed app), prepare input and run:
  `sbatch job` which contains `mpirun app (most likely: 1 core/task)`
- Can HPC performance be improved?
  Need to understand the resource usage (**Yes, it's complicated**):
  - Know the Hardware (HW) Architecture
  - Evaluate resource usage (htop, remora, etc.)
  - Know how to control/map application tasks to resources.
- Specifying where tasks execute through **affinity** (settings) can enhance performance.

- So,... it may be necessary to peek inside a node to see what is happening!

*SMT=>Hyper-Threading Technology(intel)/Cluster Multi-Threading(AMD)/SMT(IBM)

# A peak at Hardware (TACC Frontera Node)

```
$ lscpu
```

```
CPU(s):                  56
On-line CPU(s) list:     0-55
Thread(s) per core:      1
Core(s) per socket:      28
Socket(s):               2
NUMA node(s):            2
Model name:              Intel(R) Xeon(R) Platinum 8280 CPU @ 2.70GHz
CPU MHz:                 3299.853
CPU max MHz:             4000.0000
CPU min MHz:             1000.0000
L1d cache:               32K
L1i cache:               32K
L2 cache:                1024K
L3 cache:                39424K
NUMA node0 CPU(s):       0,2,4,6,8,10,12,14,16,...,38,40,42,44,46,48,50,52,54
NUMA node1 CPU(s):       1,3,5,7,9,11,13,15,17,...,39,41,43,45,47,49,51,53,55
```

Some information not shown.

# Frontera Node -- cores and memory

2-sockets, 56 Cores, Non-Uniform Memory Access (NUMA), No Hyperthreading

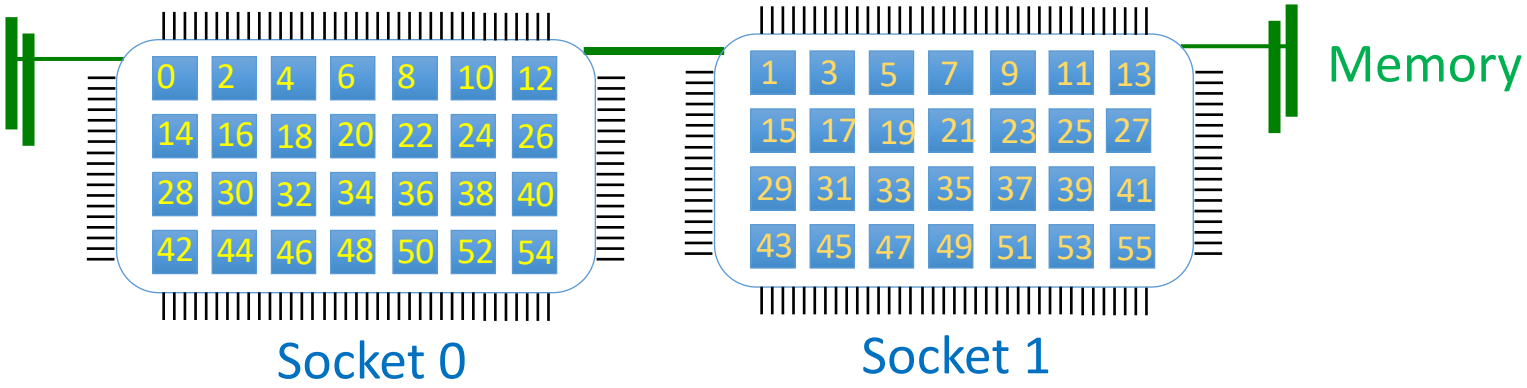

# A peak at Hardware (Lumi/G)

```
$ lscpu
```

```
CPU(s):                          128
On-line CPU(s) list:             0-127
Model name:                      AMD EPYC 7A53 64-Core Processor
Thread(s) per core:              2
Core(s) per socket:              64
Socket(s):                       1
CPU max MHz:                     2000.0000
CPU min MHz:                     1500.0000
L1d cache:                       2 MiB (64 instances)
L1i cache:                       2 MiB (64 instances)
L2 cache:                        32 MiB (64 instances)
L3 cache:                        256 MiB (8 instances)
UMA node(s):                     4
NUMA node0 CPU(s):               0-15,64-79
NUMA node1 CPU(s):               16-31,80-95
NUMA node2 CPU(s):               32-47,96-111
NUMA node3 CPU(s):               48-63,112-127
```

Some information not shown.

# A peak at Hardware (Lumi/C)

```
$ lscpu
```

```
CPU(s):                          256
On-line CPU(s) list:             0-255
Model name:                      AMD EPYC 7763 64-Core Processor
Thread(s) per core:              2
Core(s) per socket:              64
Socket(s):                       2
CPU max MHz:                     2450.0000
CPU min MHz:                     1500.0000
L1d cache:                       4 MiB (128 instances)
L1i cache:                       4 MiB (128 instances)
L2 cache:                        64 MiB (128 instances)
L3 cache:                        512 MiB (16 instances)
NUMA node(s):                    8
NUMA node0 CPU(s):               0-15,128-143
...
NUMA node6 CPU(s):               96-111,224-239
NUMA node7 CPU(s):               112-127,240-255
```

Some information not shown.

# LUMI Node -- cores and memory

2-sockets, 128 Cores, Non-Uniform Memory Access (NUMA), Hyperthreading

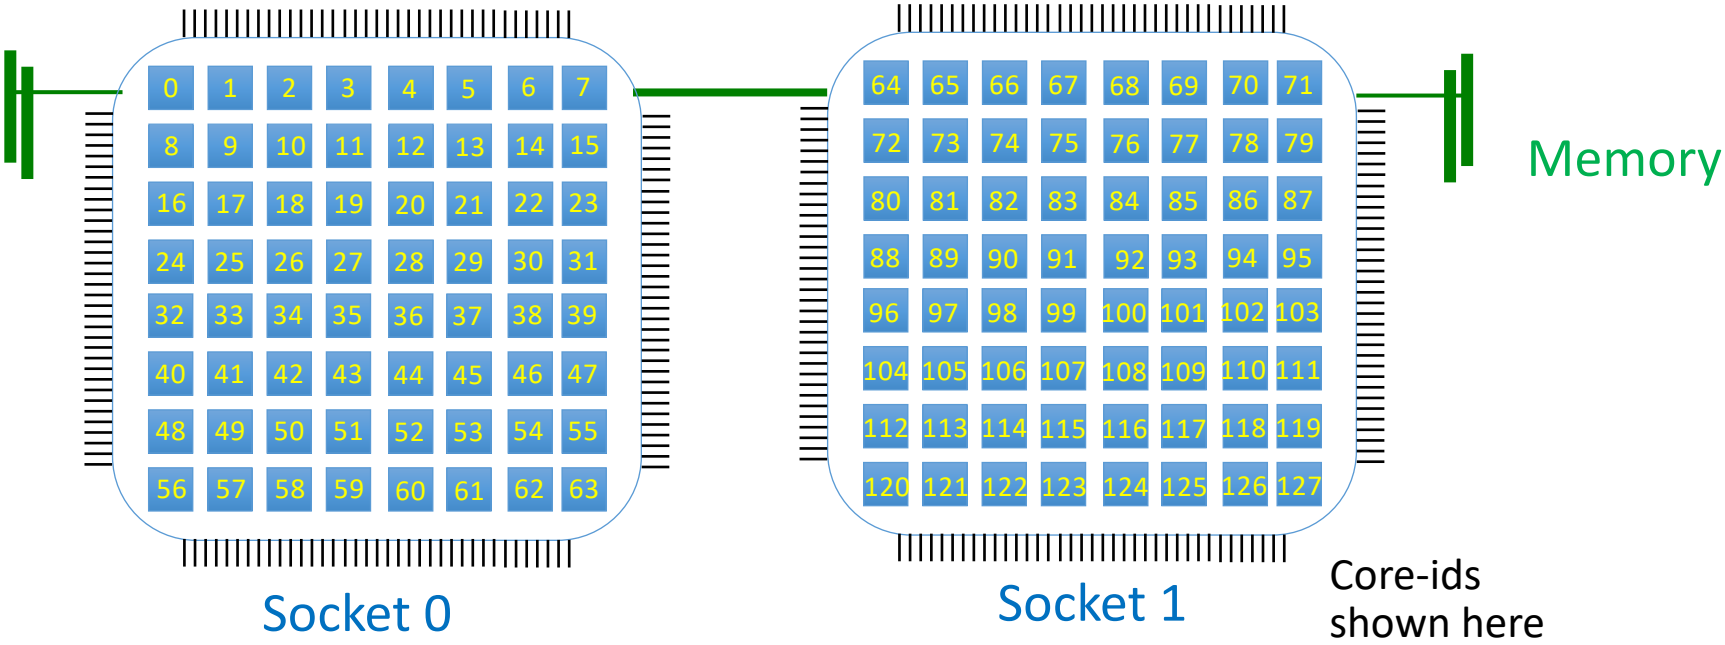

# Example Architecture

Single Socket (not always simple)

Hyper-Threading (laptop, 2) (KNL, 4 )

2 Cores (a Tile) share a 1MB L2 Cache.

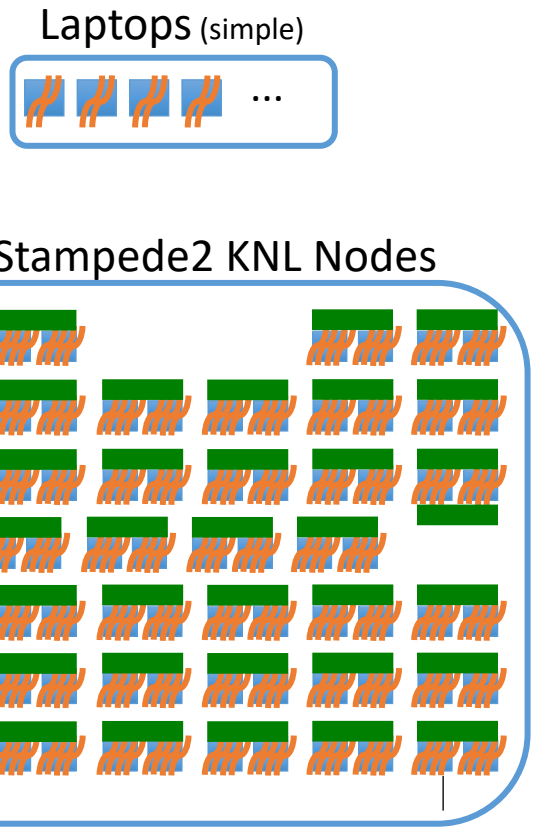

# Unix Architecture from *lscpu*

```
sp$ lscpu | grep -i 'core\|thread\|socket'
Thread(s) per core:    1
Core(s) per socket:    8
Socket(s):             2
```

```
ls5%  lscpu | grep -i 'core\|thread\|socket'
Thread(s) per core:    2
Core(s) per socket:    12
Socket(s):             2
```

```
knl$ lscpu | grep -i 'core\|thread\|socket'
Thread(s) per core:    4
Core(s) per socket:    68
Socket(s):             1
```

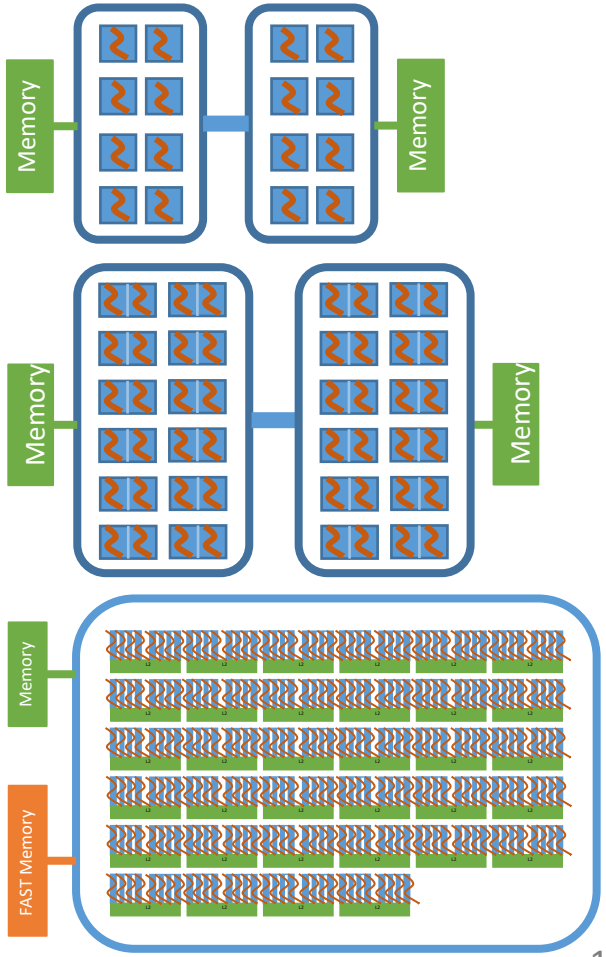


But more details can be obtained from hwloc.

# Notes from *lscpu* architectural review:

Memory:

On multi-socket nodes it might be best to bind a process to a single socket  to preserve local memory access.

Binding to a single core (when #processes <#cores) may reduce process migration and hence cache refreshes.

# When should affinity be considered?

**Affinity: Determines where application processes can/will execute.**

**Affinity: Needed for these cases:**

- processor count > process count
- processors have local and remote shared memory
- want cache to remain hot through process execution
- multiple IO clients may need separation/positioning

# Motivation

- "Often" applications run fine when using the "normal" number of processes (one process per core)
- 1st Approach to Affinity: Do Nothing, Generally Defaults are "good"
  2nd match process needs to architecture features/resources
  3rd "pin" processes for limited mobility

Need Tools to Assess Core Occupation on many-core systems:

Stampede2: KNL -- 68 cores ( 272 "logical" processors), TILES,...
Lonestar6: Milan -- 128 cores ( 128 "logical" processors), 2 SOCKETS,...
Frontera: CLK -- 56 cores ( 56 "logical" processors), 2 SOCKETS,...
LUMI-C: EPYC -- 128 cores ( 256 "logical" processors), 2 SOCKETS,...

# Outline

- Motivation
- Affinity →
  - affinity fundamentals, the affinity mask, and how it is used
- OpenMP Affinity: PROC_BIND and PLACES
- Showing Mask with *amask* utility

# CPU Affinity -- the mask

processes map onto processors


MPI rank-num or thread-num

proc-id (core-id)

A kernel bit mask (array of bits) exists for each process

# of bits = # of processors (proc-ids) set bit→ process can run on proc-id (core-id).

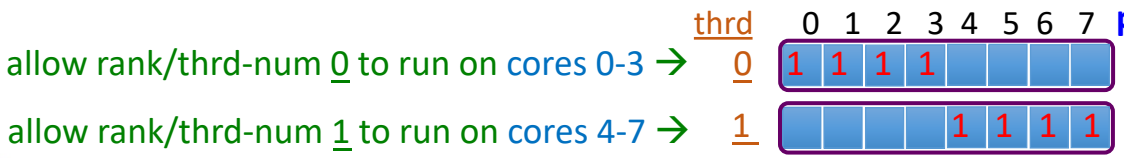


How are the proc-ids (numbers) assigned to the hardware?

# Hardware – component enumeration

- You may know how many, ... but you cannot assume anything about the numbering, even on an “identical” systems.

- Things may be according to Vendor-OEM BIOS specification, BIOS switch, Site Management, OS level, OpenMP runtime for compiler,...

“Hardware organization is unpredictable”

# e.g. setting socket affinity with proc-ids

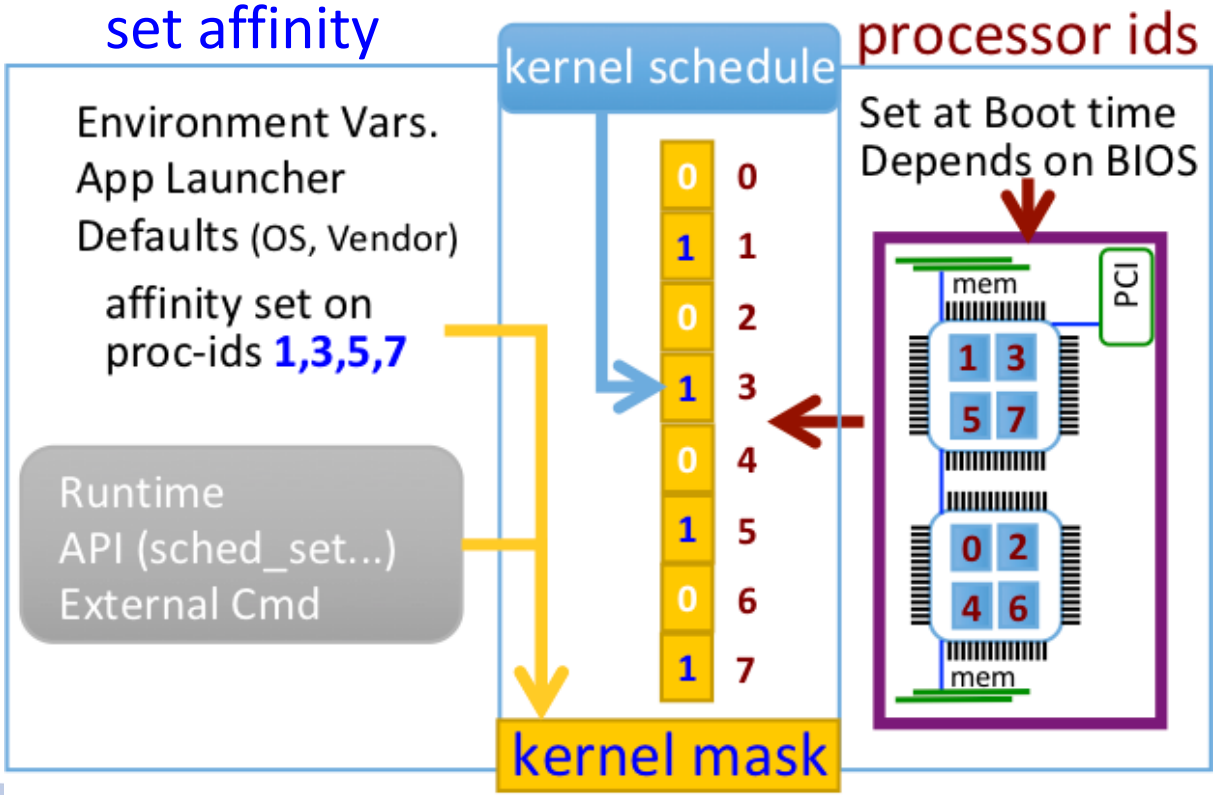


"Even-Odd Sequence"

# e.g. setting socket affinity with proc-ids

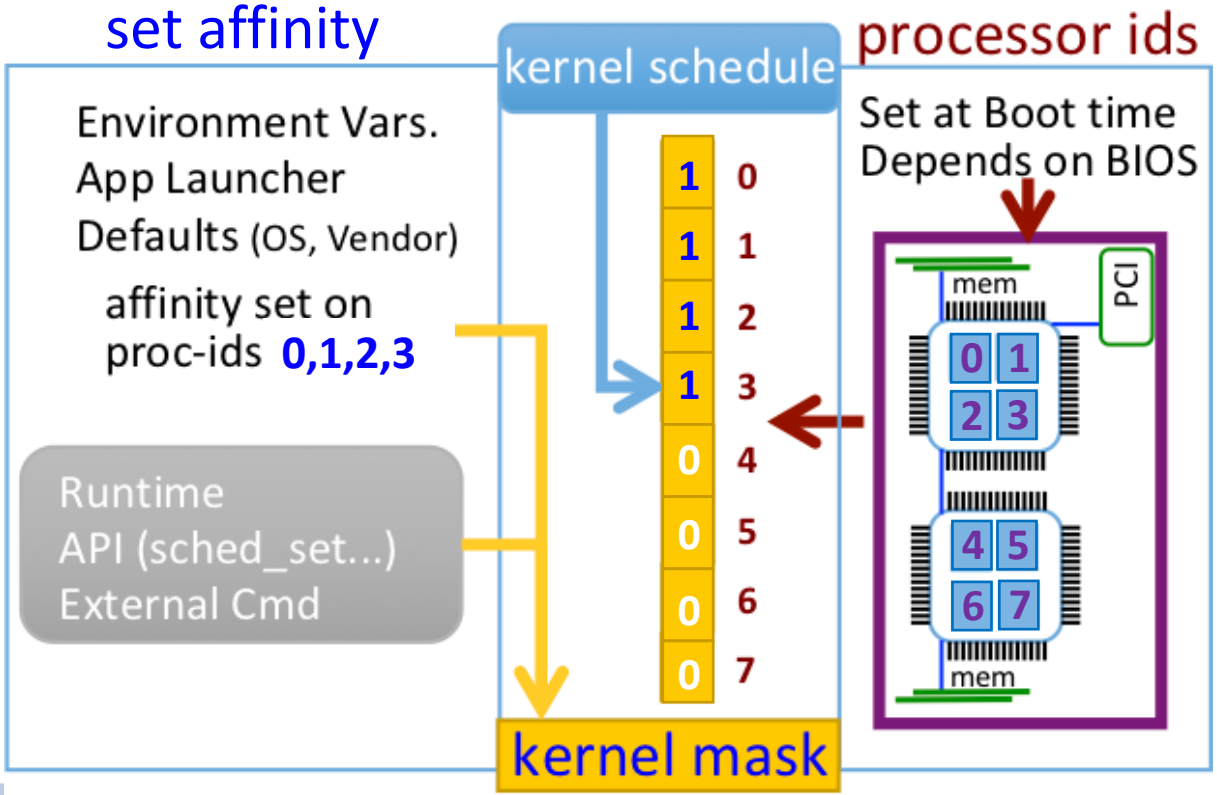


"Even-Odd Sequence"

# Setting Affinity

hwloc (lstopo) and /proc/cpuinfo are your friends

For NUMA information, see also numactl -H

# Unix hwloc utility

```
$ lstopo
```

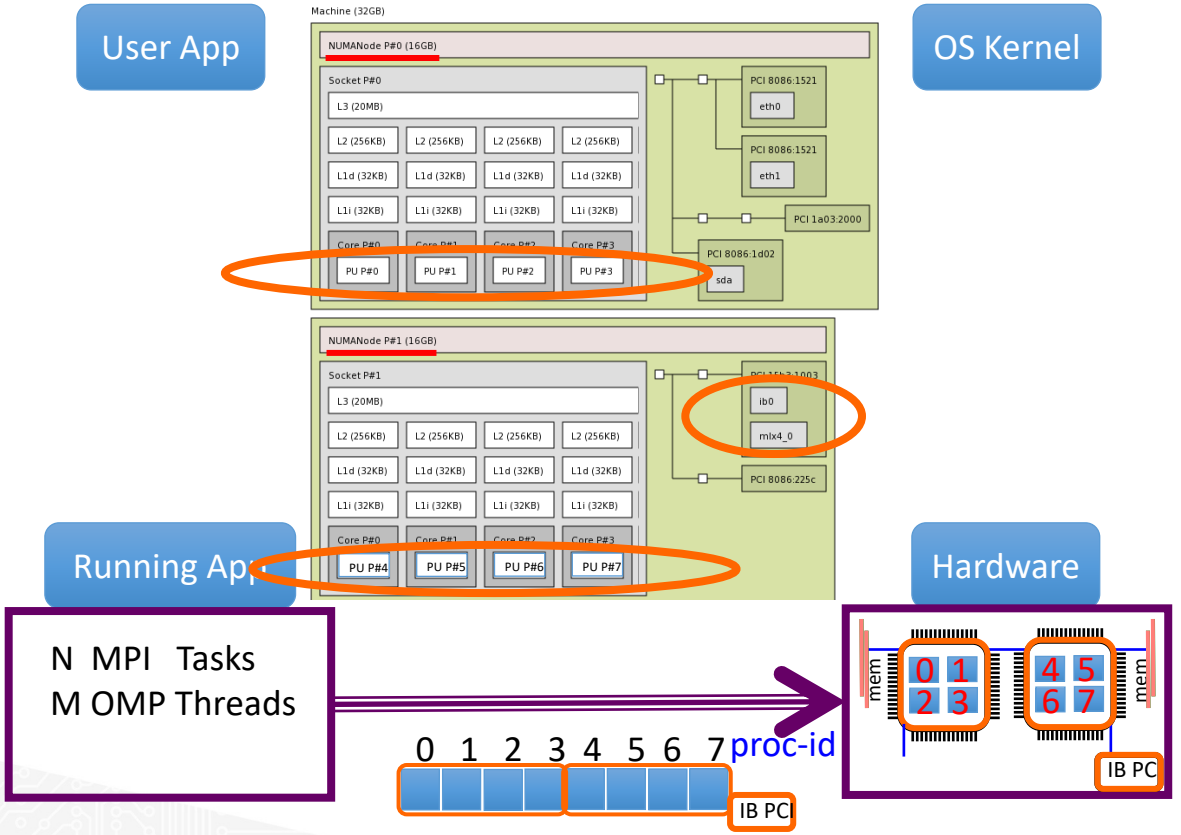


"Counting Sequence"

# Setting Affinity

## Tasks can be pinned to sets of cores or NUMA nodes

- Each task can have a "mask" of cores it may run on
- A task will can only float to allowed cores (of its mask)
- Many tools for controlling affinity:
  - numactl
  - taskset
  - sched_setaffinity (API, inside code)
  - `OMP_...` (`OpenMP`) & `GOMP_...` (GNU) & `KMP_...` (Intel)
  - MPI implementations have their own affinity tools

# Assigning threads

- How are threads assigned to proc-ids?
  - Build a list of "places" (proc-ids), or use the implementation defaults.
  - Each sequential thread number is assigned to a place.

```
#omp parallel num_threads(8)
  printf("%d\n", omp_get_thread_num());
```

A program creates 8 threads sequentially numbered

0
1
2
3
4
5
6
7

Thread numbers are assigned sequentially to places in place list.

Default place list
usually follows sequential numbering from socket to socket.
(found in `/proc/cpuinfo`)

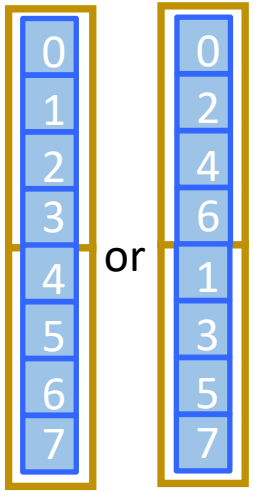


Counting Sequential  Even-Odd Sequence

# Unix Architecture from *lscpu*

```
frontera$ lscpu | grep -i 'core\|thread\|Socket'
Thread(s) per core:    1
Core(s) per socket:    28
Socket(s):             2
```

Default Place list: Known to have
Even proc-ids on socket 0.
Odd proc-ids on socket 1.

## Frontera Compute Node

One would expect the place list to be:

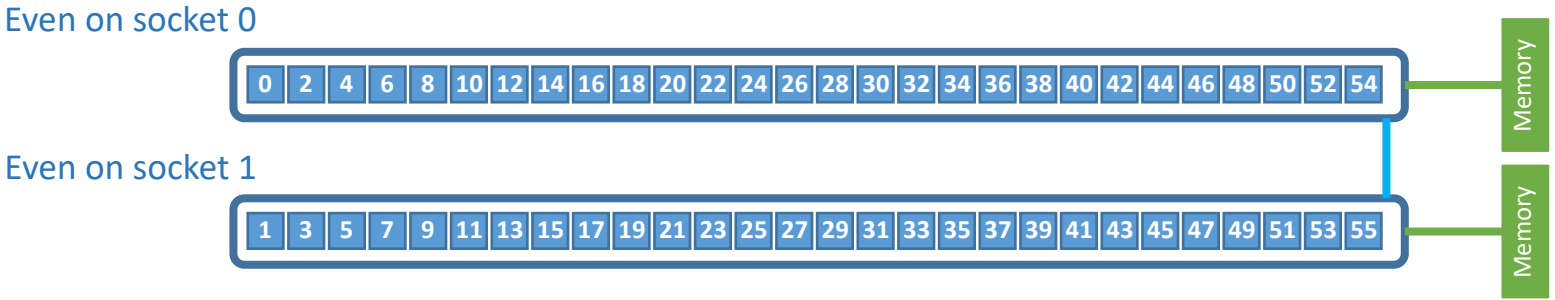

# proc-id #'s on Frontera

socket core HW-thread

```
clx% cat /proc/cpuinfo | awk ...
```

```
physical id : 0    core id  :  0    processor :  0
physical id : 0    core id  :  1    processor :  4
physical id : 0    core id  :  2    processor :  8
physical id : 0    core id  :  3    processor : 12
physical id : 0    core id  :  4    processor : 10
physical id : 0    core id  :  5    processor :  6
physical id : 0    core id  :  6    processor :  2
physical id : 0    core id  :  8    processor : 16
physical id : 0    core id  :  9    processor : 20
physical id : 0    core id  : 10    processor : 24
physical id : 0    core id  : 11    processor : 26
physical id : 0    core id  : 12    processor : 22
physical id : 0    core id  : 13    processor : 18
physical id : 0    core id  : 14    processor : 14
physical id : 0    core id  : 16    processor : 28
physical id : 0    core id  : 17    processor : 32
physical id : 0    core id  : 18    processor : 36
physical id : 0    core id  : 19    processor : 40
physical id : 0    core id  : 20    processor : 38
physical id : 0    core id  : 21    processor : 34
physical id : 0    core id  : 22    processor : 30
physical id : 0    core id  : 24    processor : 44
physical id : 0    core id  : 25    processor : 48
physical id : 0    core id  : 26    processor : 52
physical id : 0    core id  : 27    processor : 54
physical id : 0    core id  : 28    processor : 50
physical id : 0    core id  : 29    processor : 46
physical id : 0    core id  : 30    processor : 42
...
```

Note the non-sequential processor (proc-id) numbering.

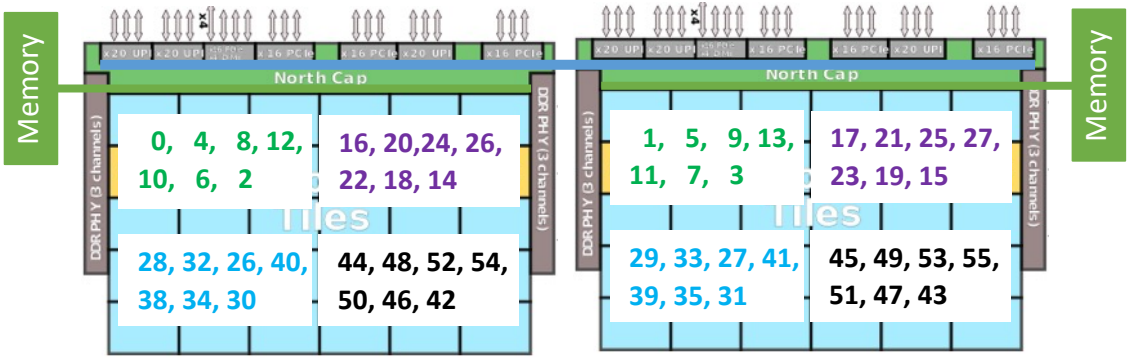


```
#pragma omp parallel
  printf("%d\n", omp_get_thead_num());
```

Thread numbers are assigned sequentially to core-id sequence in /proc/cpuinfo.

# lstopo on Frontera node

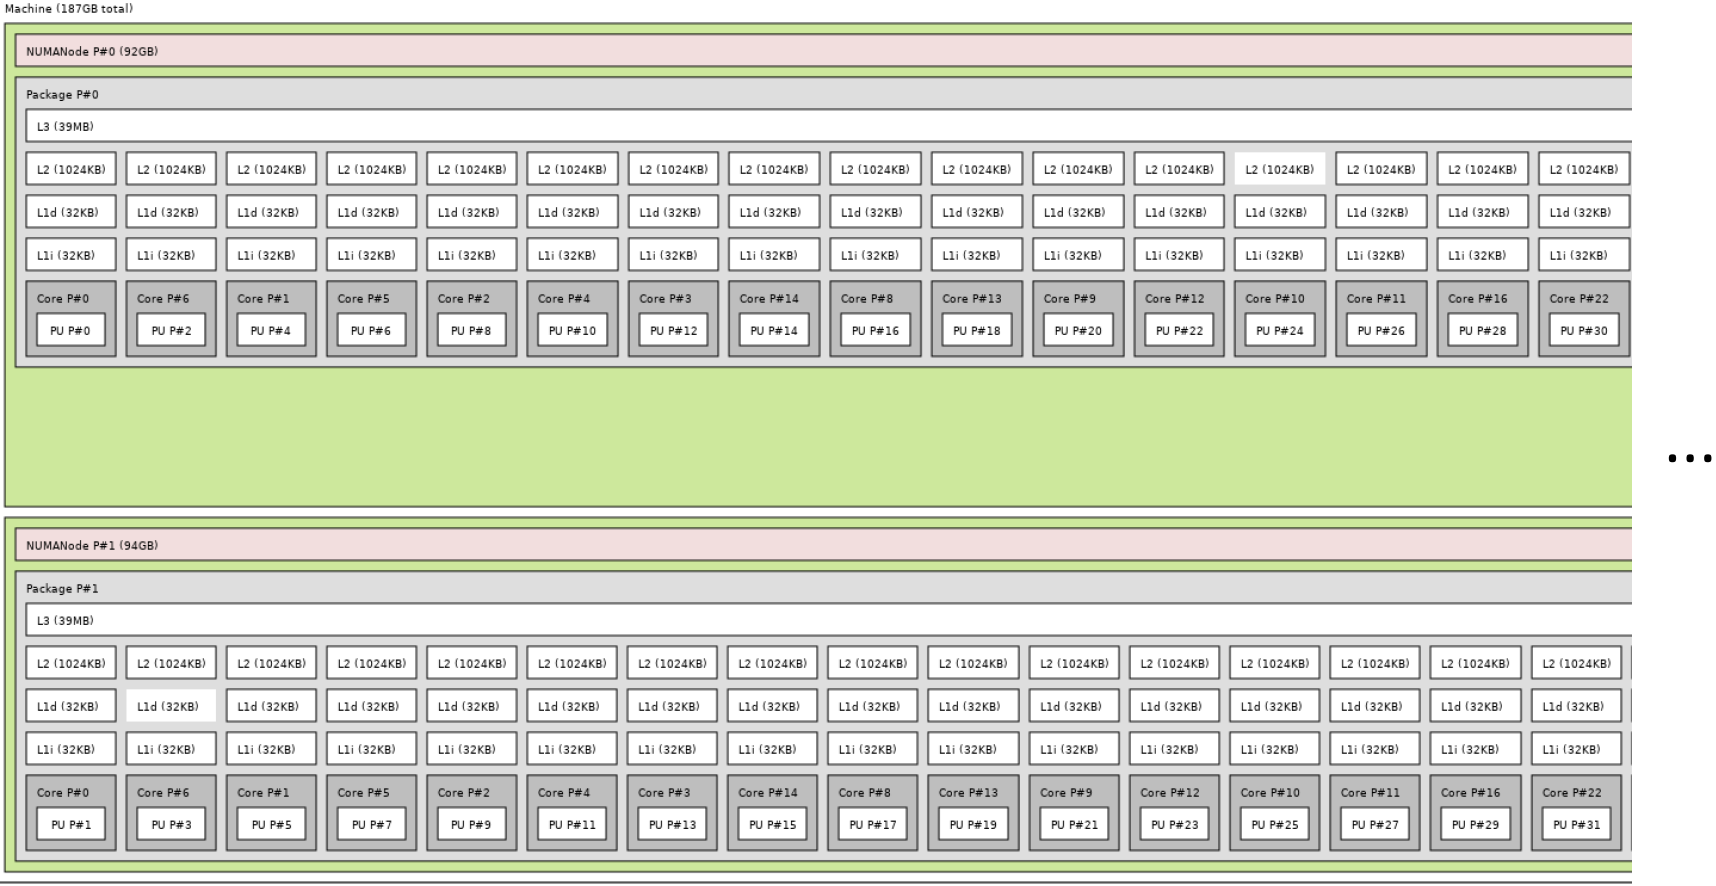


...

# lstopo on Frontera node (cont. 1)

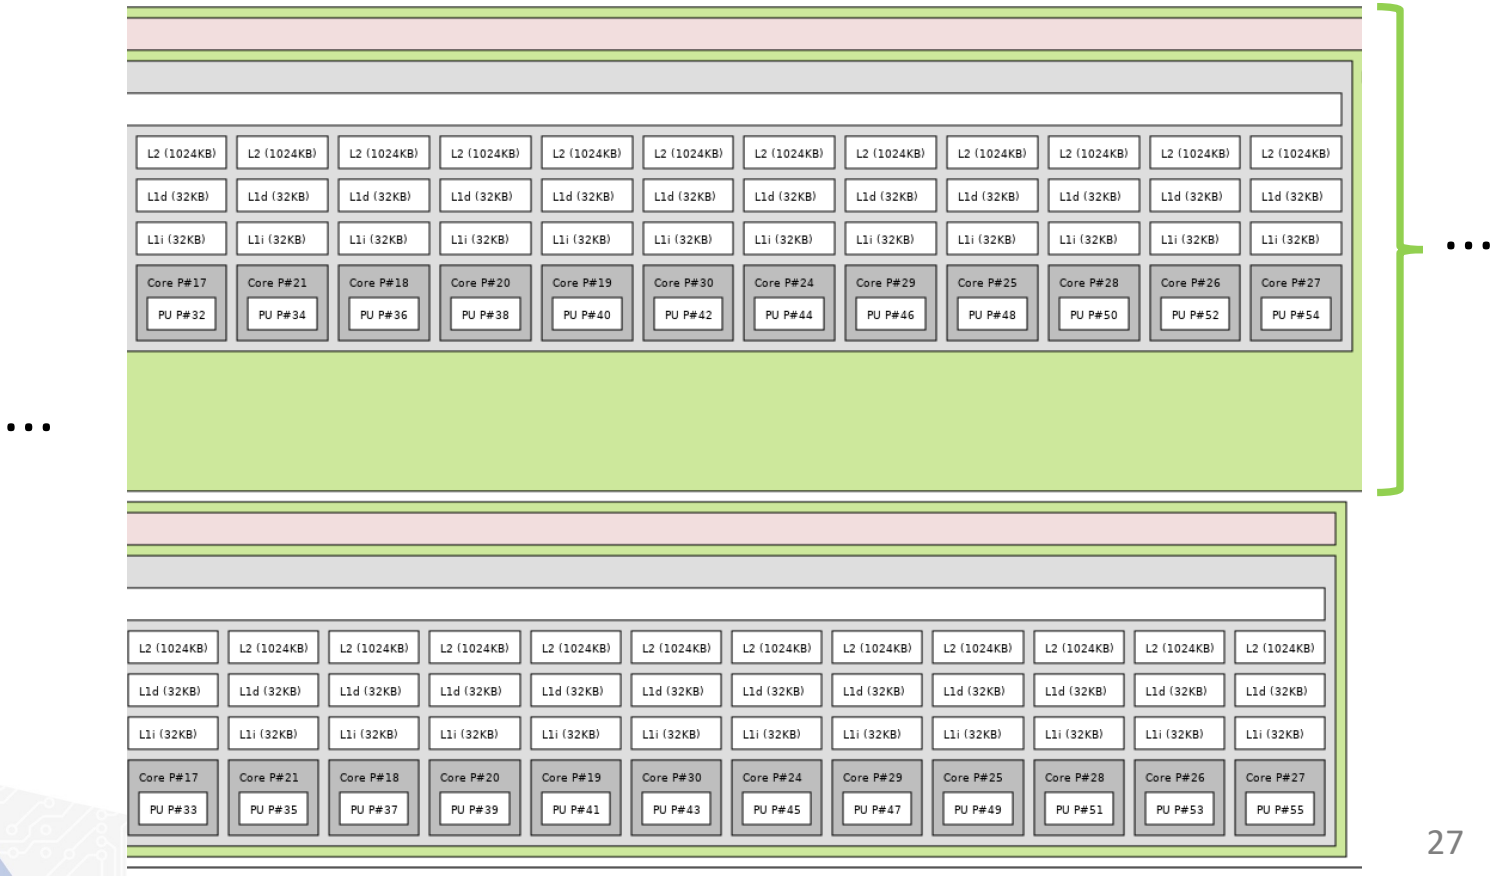

# Istopo on Frontera node (cont. 2)

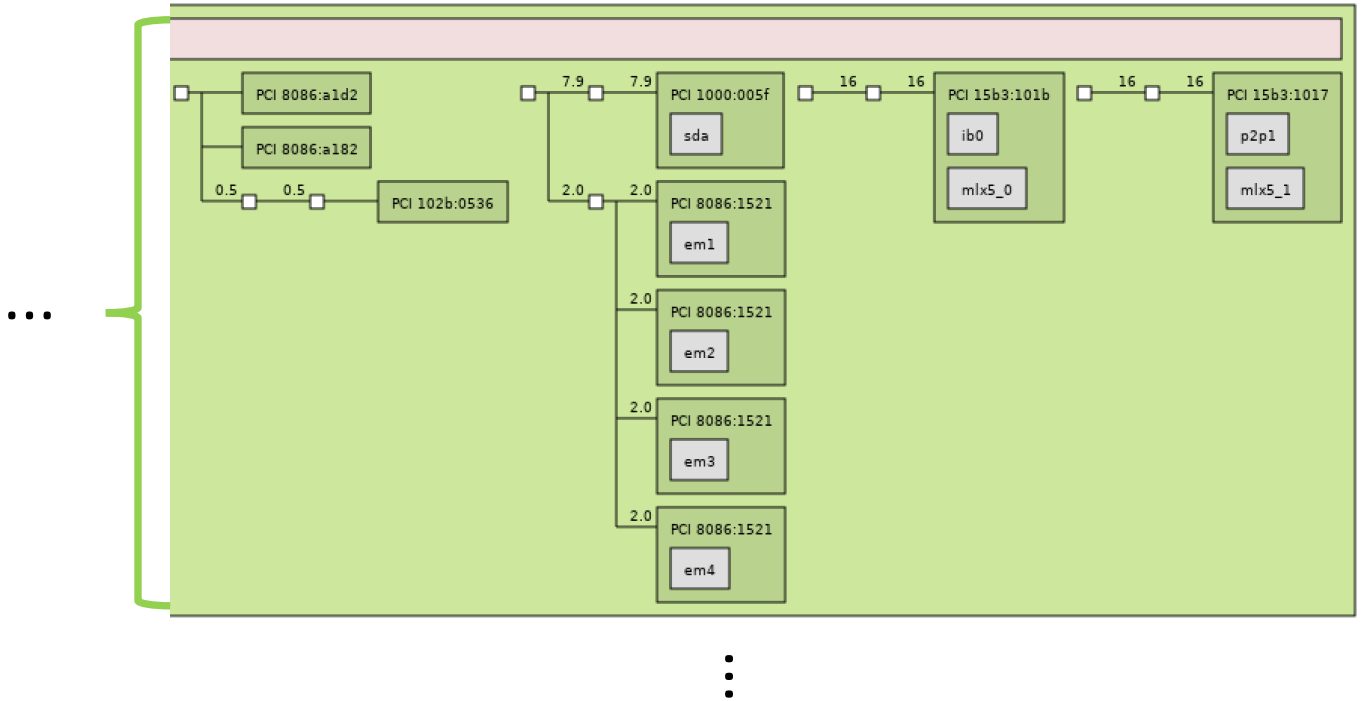

# Unix Architecture from *lscpu*

## Frontera Compute Node

Think of thread assignment ordering this way.

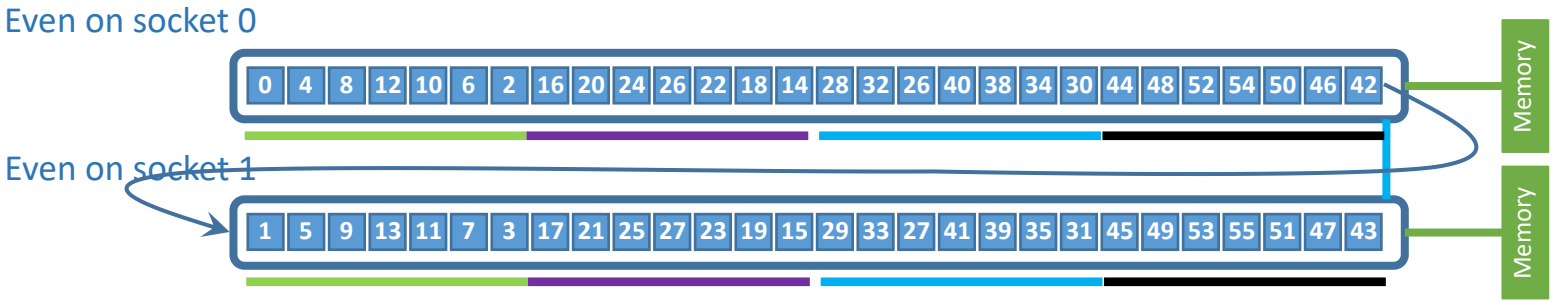

# NUMA nodes on AMD EPYC

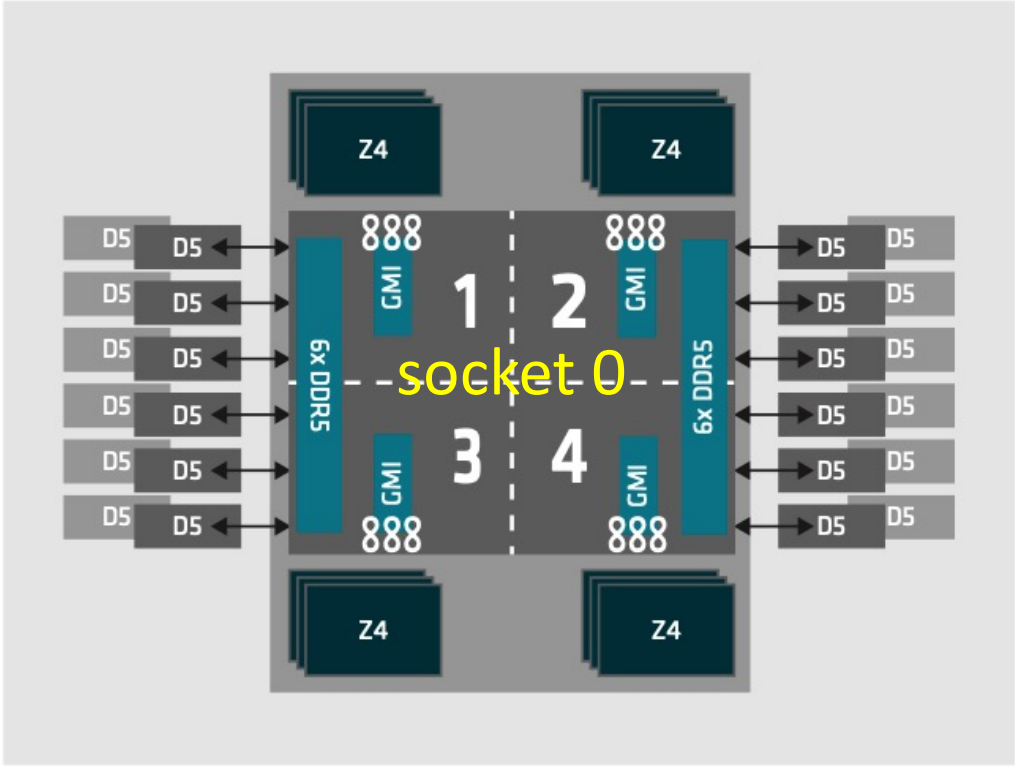


*Figure 7: Dividing the AMD EPYC processor into four NUMA domains can give small performance improvements for some applications*

https://www.amd.com/system/files/documents/4th-gen-epyc-processor-architecture-white-paper.pdf

# LUNI/C Architecture from *lstopo and numactl -H*

```
$ lstopo
```

```
Machine (251GB total)
  Package L#0
    Group0 L#0
      NUMANode L#0 (P#0 31GB)
      L3 L#0 (32MB)
        L2 L#0 (512KB) + L1d L#0 (32KB) + L1i L#0 (32KB) + Core L#0
          PU L#0 (P#0)
          PU L#1 (P#128)
        L2 L#1 (512KB) + L1d L#1 (32KB) + L1i L#1 (32KB) + Core L#1
          PU L#2 (P#1)
          PU L#3 (P#129)
        ...
        L2 L#15 (512KB) + L1d L#15 (32KB) + L1i L#15 (32KB) + Core L#15
          PU L#30 (P#15)
          PU L#31 (P#143)
    Group0 L#1
      NUMANode L#1 (P#1 31GB)
      ...
      HostBridge
        PCIBridge
          PCI 41:00.0 (Ethernet)
            Net "nmn0"


    Group0 L#2
      NUMANode L#2 (P#2 31GB)
      L3 L#4 (32MB)
        L2 L#32 (512KB) + L1d L#32 (32KB) + L1i L#32 (32KB) + Core L#32
          PU L#64 (P#32)
          PU L#65 (P#160)
```

# Lumi C Architecture

## LUMI/C Node

Sequential core-id numbers from socket 0 and continuing on socket1

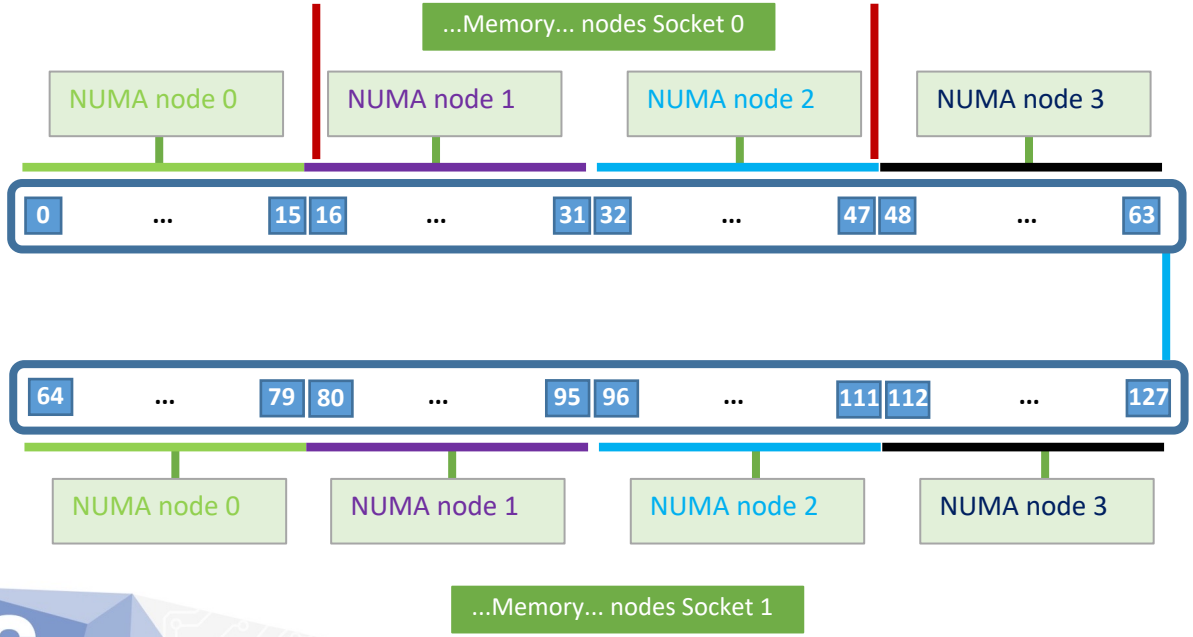

# Outline

- Motivation
- Affinity →
  - affinity fundamentals, the affinity mask, and how it is used
- OpenMP Affinity: **PROC_BIND and PLACES**
- Showing Mask with *amask* utility

# OpenMP Affinity

Portable: works with GNU, Intel, IBM, Oracle, …

There are two components to setting affinity:

(after setting the number of threads)

**Distribution Policy**: PROC_BIND policy

**Set Locations:** PLACES User list of proc-ids or abstract set

List sequence by by core-id order in /proc/cpuinfo.

Environment Variables:

```
export PROC_BIND=<...> PLACES=<...>
```

# OpenMP Affinity (Policy)

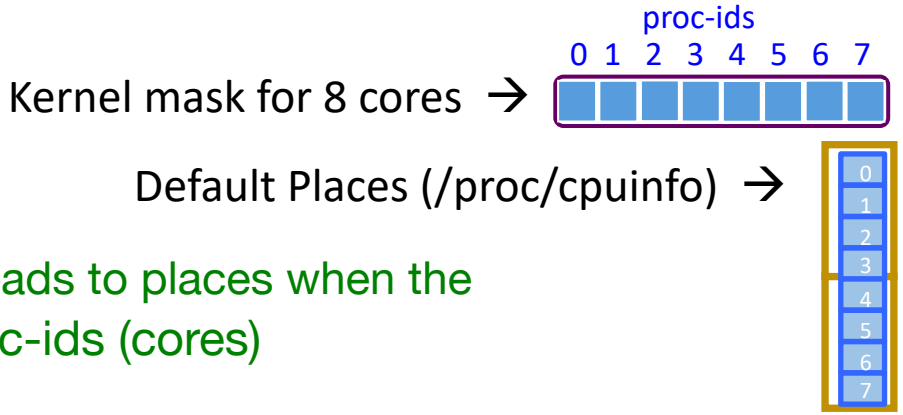


- proc_bind: describes how to distribute threads to places when the number of threads is less than number of proc-ids (cores)

- export OMP_NUM_THREADS=4 # (8-core node)
- export OMP_PROC_BIND=close

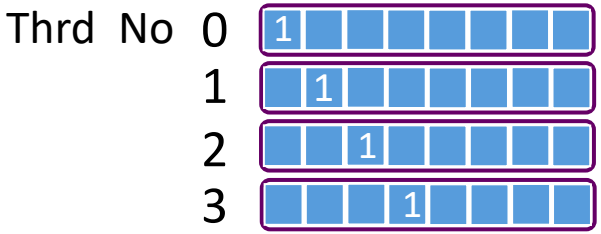


- export OMP_NUM_THREADS=4 # (8-core node)
- export OMP_PROC_BIND=spread

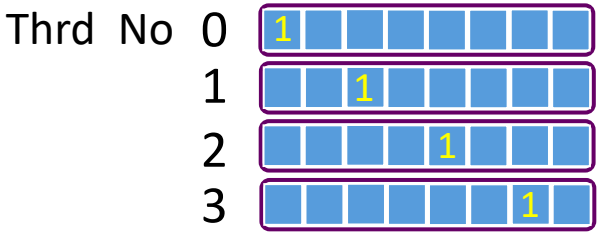

# OpenMP Affinity (User Place List)

Kernel mask for 8 cores →

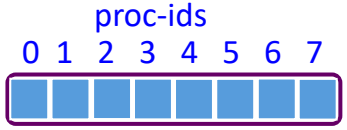


- place: a set of proc-ids* where a thread can execute

{4, 5, 6, 7}


OMP_NUM_THREADS=3

- place list: list of places assigned to threads in order (for thread no. 0, 1, 2, ...)

{0, 2}, {7, 6} , {3, 4, 5}

Thrd no.

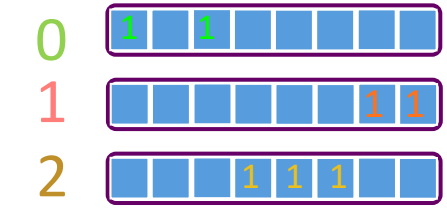


*place number = "smallest unit of execution"= proc-id # in Linux systems

# OpenMP Affinity (More on Places)

- Places List
  - export OMP_PLACES=<place_list>

    A place: {0}
    Place List: {0},{1},{2},{3}
    Interval notation: {<place>}:<len>:<stride> e.g. {0}:4:1

    List of proc-ids

  - export OMP_PLACES=<abstract_name>

    Abstract name: sockets, cores*, threads

    Defines a set of places.

    (*cores* default on Frontera)

The precise definitions of the abstract names are implementation defined. An implementation may also add abstract names as appropriate for the target platform.

# OpenMP Affinity (Place assignments)

- In a Parallel region thread numbers 0,...,N-1 are assigned sequentially to places in the place list.

```
#pragma omp parallel
printf("%d\n",omp_get_thread_num());
```

```
OMP_PLACES=\
"{20},{10},{12},{14,16}"
```

```
0 ⟷ 20
1 ⟷ 10
2 ⟷ 12
3 ⟷ 14,16
```

Frontera

```
OMP_PLACES=cores
OMP_PROC_BIND=close
(uses sequence in
/proc/cpuinfo)
```

```
0
4
8
12
```

[Frontera OMP_PROC_PLACES default is cores]

# OpenMP Affinity
# PROC_BIND Distribution

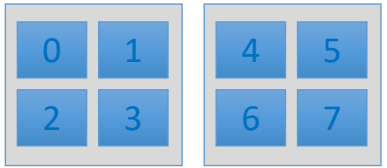

2 sockets x 4 cores

```
export OMP_NUM_THREADS=4
export OMP_PROC_BIND=close

!$omp parallel private(thrd_num);
thrd_num=omp_get_thread_num();
```

```
export OMP_NUM_THREADS=4
export OMP_PROC_BIND=spread

!$omp parallel private(thrd_num);
thrd_num=omp_get_thread_num();
```


COMPACT PACKING


SCATTER PACKING

# OpenMP Affinity Places

```
export OMP_NUM_THREADS=4
export OMP_PLACES='{0},{1},{2},{3}'

!$omp parallel private(thrd_num);
thrd_num=omp_get_thread_num();
```

```
export OMP_NUM_THREADS=4
export OMP_PLACES='{0},{2},{4},{6}'

!$omp parallel private(thrd_num);
thrd_num=omp_get_thread_num();
```


COMPACT PACKING

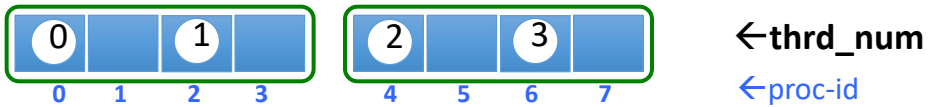


SCATTER PACKING

# OpenMP Affinity
# Places → Interval Expression

```
export OMP_NUM_THREADS=4
export OMP_PLACES='{0}:4'

!$omp parallel private(thrd_num);
thrd_num =omp_get_thread_num();
```

```
export OMP_NUM_THREADS=4
export OMP_PLACES='{0}:4:2'

!$omp parallel private(thrd_num);
thrd_num =omp_get_thread_num();
```


COMPACT PACKING

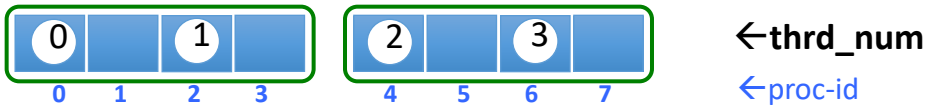


SCATTER PACKING

# OpenMP Affinity
# Places → Interval Expression

```
export OMP_NUM_THREADS=2
export OMP_PLACES='{0,1}:2:2'

!$omp parallel private(thrd_num);
thrd_num =omp_get_thread_num();
```

```
export OMP_NUM_THREADS=2
export OMP_PLACES='{0,1}:2:4'

!$omp parallel private(thrd_num);
thrd_num =omp_get_thread_num();
```


COMPACT PACKING


SCATTER PACKING

# OpenMP Affinity
# Places → abstract name

## Hyper-Threading Enabled

```
export OMP_NUM_THREADS=8
export OMP_PLACES=threads

!$omp parallel private(thrd_num);
thrd_num =omp_get_thread_num();
```

```
export OMP_NUM_THREADS=8
export OMP_PLACES=cores

!$omp parallel private(thrd_num);
thrd_num =omp_get_thread_num();
```


Binding to Single HW-thread


Binding to Core

←thrd_num
←proc-id

HW-thread

# OpenMP Affinity
# Places expression

## Hyper-Threading Enabled

```
export OMP_NUM_THREADS=8
export OMP_PLACES='{0}:8'   CASE 1
export OMP_PLACES='{8}:8'   CASE 2

!$omp parallel private(thrd_num);
thrd_num =omp_get_thread_num();
```

```
export OMP_NUM_THREADS=8
export OMP_PLACES='{0,8}:8'


!$omp parallel private(thrd_num);
thrd_num =omp_get_thread_num();
```


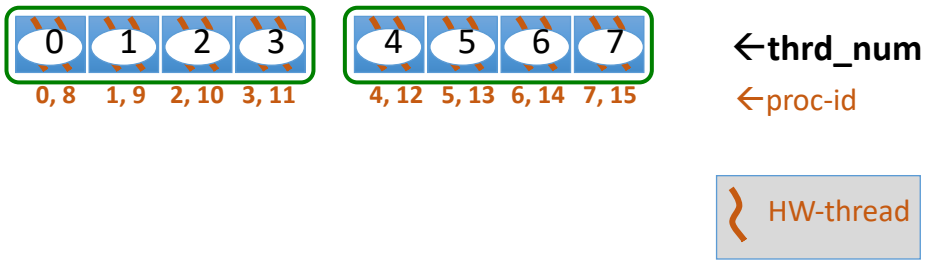


Binding to Single HW-thread

Binding to Core

# OpenMP Affinity
# Places abstract name

Frontera

even proc-ids: 0,4,8,12,10,6,2,16...

odd proc-ids: 1,5,9,13,11,7,3,17...

```
export OMP_NUM_THREADS=8
export OMP_PLACES=cores
export OMP_PROC_BIND=close

!$omp parallel private(thrd_num);
thrd_num=omp_get_thread_num();
```

```
export OMP_NUM_THREADS=8
export OMP_PLACES=cores
export OMP_PROC_BIND=spread

!$omp parallel private(thrd_num);
thrd_num=omp_get_thread_num();
```

←thrd_num

# OpenMP Affinity Places List

Frontera

even proc-ids: 0,4,8,12,10,6,2,16...

odd proc-ids: 1,5,9,13,11,7,3,17...

evens on socket 1

```
export OMP_NUM_THREADS=4
export OMP_PLACES='{0,2,4,6,8,10,12}:4:14'

!$omp parallel private(thrd_num);
thrd_num=omp_get_thread_num();
```

odds on socket 1

```
export OMP_NUM_THREADS=4
export OMP_PLACES='{1,3,5,7,9,11,13}:4:14'

!$omp parallel private(thrd_num);
thrd_num=omp_get_thread_num();
```


even proc-ids 0,2,4,...,54

odd proc-ids 1,3,5,...,55

←thrd_num

←proc-id

# Outline

- Motivation
- Affinity →
  - affinity fundamentals, the affinity mask, and how it is used
- OpenMP Affinity: PROC_BIND and PLACES
- Showing Mask with *amask* utility

# amask


- Report for an affinity env.
  - Executable command: amask_<*type*>  types:{omp, mpi, hybrid}.
  - Can instrument code:  amask_<*type*>() argumentless function calls.

Syntax:

amask_<*type*>  [-h]  [-w#]

help
wait #sec with load

Usage:  module load amask

| | | |
|---|---|---|
| export OMP_<affinity>= … | amask_omp | -- for pure OpenMP |
| export I_MPI_<affinity>= … | mpirun amask_mpi | -- for pure MPI |
| export OMP_<affinity>= … \<br>I_MPI_<affinity>= … | mpirun amask_hybrid | -- MPI + OpenMP |

# Viewing Affinity mask with amask

Old Stampede: 16-core

export OMP_NUM_THREADS=8 OMP_PROC_BIND=spread

amask_omp

```
        0   proc-id    15
thrd    v              v
 | 0   1000000000000000
 v 1   0010000000000000
   2   0000100000000000
   3   0000001000000000
   4   0000000010000000
   5   0000000000100000
   6   0000000000001000
   7   0000000000000010
```

Bit Mask

```
                proc-id
thrd    |  0       |    10    |
   0    1---------------
   1    --1-------------
   2    ----1-----------
   3    ------1---------
   4    --------1-------
   5    ----------1-----
   6    ------------1---
   7    --------------1-
rank    |  0       |    10    |
```

More Readable

```
                proc-id
thrd    |  0       |    10    |
   0    0---------------
   1    --2-------------
   2    ----4-----------
   3    ------6---------
   4    --------8-------
   5    ----------0-----
   6    ------------2---
   7    --------------4-
rank    |  0       |    10    |
```

Even More Readable

# Frontera

$ OMP_NUM_THREADS=2 amask_omp

NO OpenMP Affinity

```
thrd |    0    |   10    |   20    |   30    |   40    |   50    |
0000 01234567890123456789012345678901234567890123456789012345
0001 01234567890123456789012345678901234567890123456789012345
```

$ OMP_NUM_THREADS=2 OMP_PROC_BIND=spread amask_omp

1 thread per socket

```
thrd |    0    |   10    |   20    |   30    |   40    |   50    |
0000 0-------------------------------------------------------
0001 -1------------------------------------------------------
```

$ OMP_NUM_THREADS=2 OMP_PROC_BIND=close amask_omp

Both threads on 1 socket

```
thrd |    0    |   10    |   20    |   30    |   40    |   50    |
0000 0-------------------------------------------------------
0001 ----4--------------------------------------------------
```

# Frontera

```
$ OMP_NUM_THREADS=28 \
  OMP_PROC_BIND=close \
  amask_omp
```

```
thrd
|    0    |   10    |   20    |   30    |   40    |   50    |
0000 0-------------------------------------------------------
0001 ----4-------------------------------------------------------
0002 --------8-------------------------------------------------------
0003 ------------2-------------------------------------------------------
0004 ----------0-------------------------------------------------------
0005 ------6-------------------------------------------------------
0006 --2-------------------------------------------------------
0007 ----------------6-------------------------------------------------------
0008 ------------------0-------------------------------------------------------
0009 ----------------------4-------------------------------------------------------
0010 ------------------------6-------------------------------------------------------
0011 --------------------2-------------------------------------------------------
0012 ------------------8-------------------------------------------------------
0013 --------------4-------------------------------------------------------
0014 --------------------------8-------------------------------------------------------
0015 ----------------------------2-------------------------------------------------------
0016 ------------------------------------6-------------------------------------------------------
0017 ----------------------------------------0-------------------------------------------------------
0018 ------------------------------------8-------------------------------------------------------
0019 ----------------------------------4-------------------------------------------------------
0020 --------------------------------0-------------------------------------------------------
0021 ----------------------------------------------4-------------------------------------------------------
0022 ------------------------------------------------8-------------------------------------------------------
0023 ----------------------------------------------------2-------------------------------------------------------
0024 ------------------------------------------------------4-------------------------------------------------------
0025 ----------------------------------------------------0-------------------------------------------------------
0026 ------------------------------------------6-------------------------------------------------------
0027 ----------------------------------------2-------------------------------------------------------
```

## /proc/cpuinfo

```
processor :  0
processor :  4
processor :  8
processor : 12
processor : 10
processor :  6
processor :  2
processor : 16
processor : 20
processor : 24
processor : 26
processor : 22
processor : 18
processor : 14
processor : 28
processor : 32
processor : 36
processor : 40
processor : 38
processor : 34
processor : 30
processor : 44
processor : 48
processor : 52
processor : 54
processor : 50
processor : 46
processor : 42
```

# Frontera

```
$ export OMP_NUM_THREADS=4
$ export OMP_PLACES='{0,2,4,6,8,10,12}:4:14'
$ amask_omp
```

All on the same socket (even proc-ids)

```
thrd |    0    |   10    |   20    |   30    |   40    |   50    |
0000 0-2-4-6-8-0-2-------------------------------------------------
0001 --------------4-6-8-0-2-4-6---------------------------------
0002 ----------------------------8-0-2-4-6-8-0-------------------
0003 ------------------------------------------2-4-6-8-0-2-4-
```

```
$ export OMP_PLACES='{0,2,4,6,8,10,12}:4:14, {1,3,5,7,9,11,13}:4:14'
$ export OMP_NUM_THREADS=8
$ amask_omp
```

Across sockets

```
thrd |    0    |   10    |   20    |   30    |   40    |   50    |
0000 0-2-4-6-8-0-2-------------------------------------------------
0001 --------------4-6-8-0-2-4-6---------------------------------
0002 ----------------------------8-0-2-4-6-8-0-------------------
0003 ------------------------------------------2-4-6-8-0-2-4-
0004 -1-3-5-7-9-1-3-----------------------------------------------
0005 --------------5-7-9-1-3-5-7---------------------------------
0006 ----------------------------9-1-3-5-7-9-1-------------------
0007 ------------------------------------------3-5-7-9-1-3-5
```

TACC

# What about … hyper-threading?

- With Hyper-Threading it is more more reasonable to report the SMTs (hardware threads) together.

# What about hyper-threading…

Will skip in short course.

Hyper-Threaded systems

2 x 12 cores → 48 hardware threads

```
$ export OMP_NUM_THREADS=4 OMP_PLACES=cores
$ amask_omp
```

defaults to “core-id” mask when hyper-threading is on

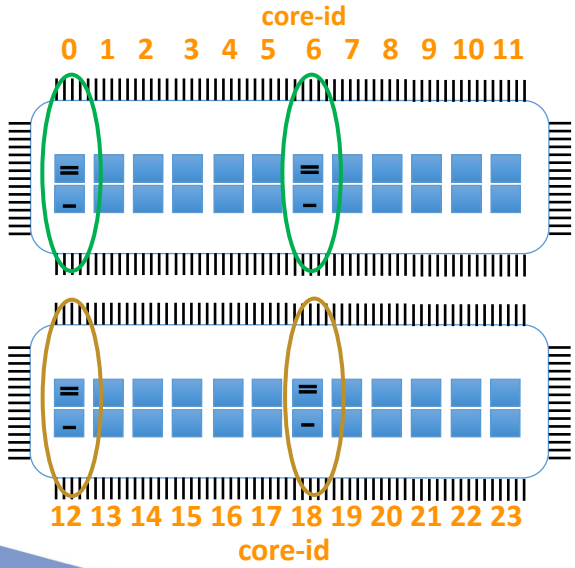


```
thrd |    0    |   10    |   20   ← core-id
0000 0=======================      HW-thread 0
     0-----------------------      HW-thread 1
0001 ======6=================
     -----6-----------------
0002 ==========2============
     ----------2------------
0003 ================8=====
     ----------------8-----
```

# proc-id #'s in /cpu/info

```
ls5%  lscpu | grep -i 'core\|thread\|socket'
Thread(s) per core:    2
Core(s) per socket:    12
Socket(s):             2
```

Hierarchy summary

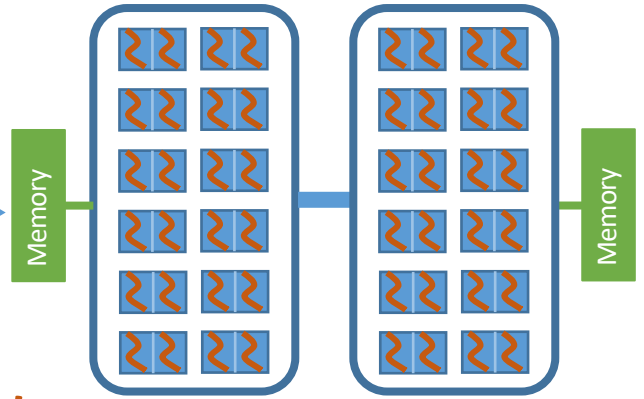


**ls5%** cat /proc/cpuinfo ...awk* ...

```
physical id : 0    core id  : 0    processor : 0
physical id : 0    core id  : 0    processor : 24
physical id : 0    core id  : 1    processor : 1
physical id : 0    core id  : 1    processor : 25
...
physical id : 1    core id  : 0    processor : 12
physical id : 1    core id  : 0    processor : 36
physical id : 1    core id  : 1    processor : 13
...
```

proc-id #s

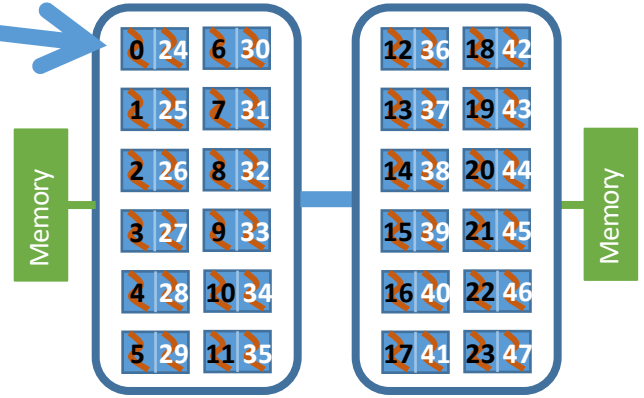


```
ls5% * awk '/processor|core id|physical id/ {arr[j++]=$0};
       END{for(i=0;i<j;i+=3) {printf "%-20s %-20s %-20s\n",
       arr[i+1],arr[i+2],arr[i]} }' /proc/cpuinfo | \
       sed -e 's/[ \t]/ /g' | sort -n -k4,4 -k8,8 -k11,11
```

# Questions!