



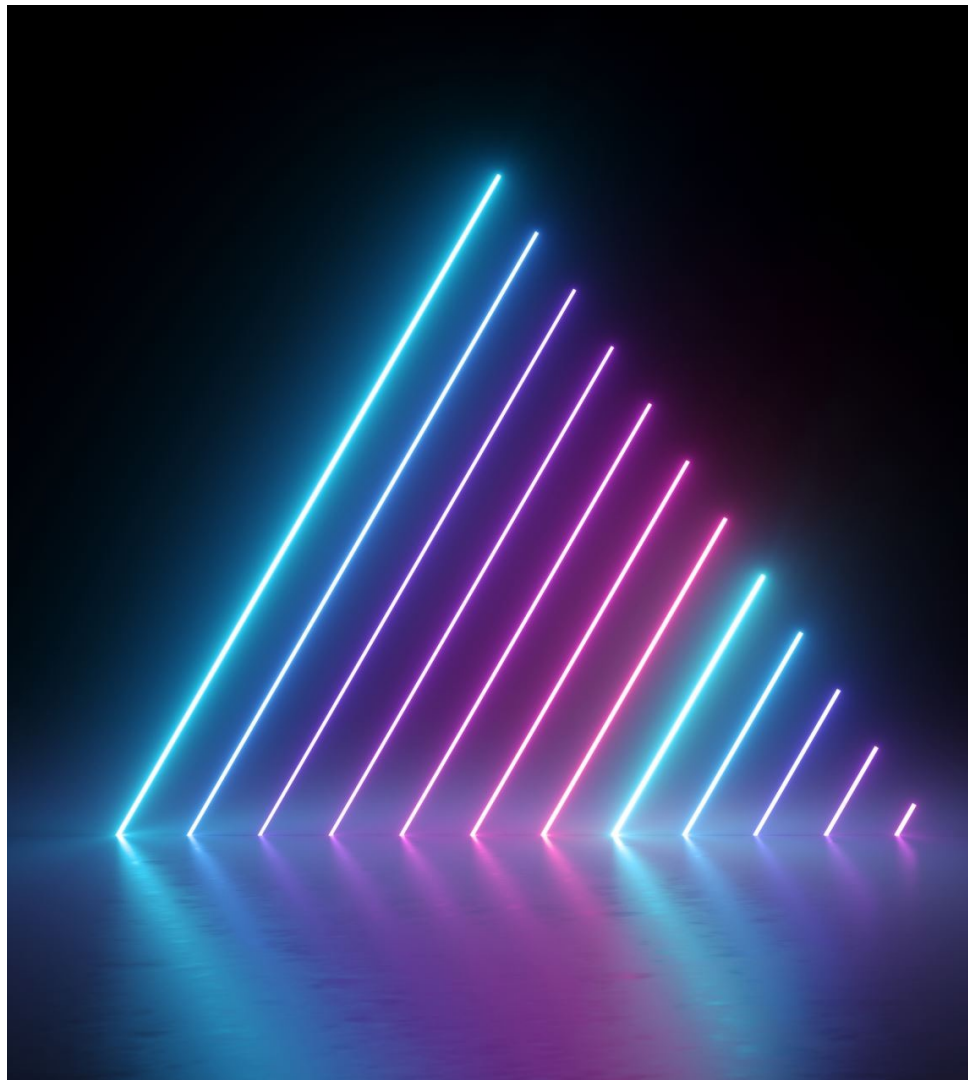
Schedule

- Beyond Threads
- Tasking
- SIMD
- Affinity
- C++
- OpenMP Offloading
 - GPUs in HPC
 - Introduction to the offloading paradigms
 - OpenMP target directive
 - OpenMP Mapping
 - OpenMP Variants, Metadata, interoperabilites

OpenMP beyond basics

CSC Summer Institute
October 9-11, 2023

Kent Milfeld (TACC)
Emanuele Vitali (CSC)



Slides available at:



General

Parallelize for orders of magnitude better performance.

CPU/GPU=> teams/grids => threads+SIMD/SMT

Teach basics with concepts, not just syntax

- This is the debut for this course – the schedule may vary
- Content is about Concepts: Describes Features, Syntax, and use with examples (+labs).
- Some content comes from other tutorials/classes/documents (sources are referenced)



Outline of Intro

- History of OpenMP
- Summary of the Basics
- Hints, tricks and and Tips for OpenMP development
- Getting setup for Lumi/Frontera



OpenMP



OpenMP beyond threads

OpenMP is NOT just about threads

OpenMP deals with multiple levels of Parallelism

target

Across Devices



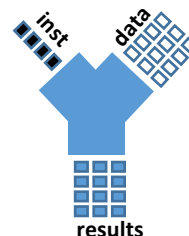
parallel do/for, task

Across Cores



simd

Across SIMD Lanes





OpenMP beyond threads

OpenMP is NOT just about threads

OpenMP deals with Memory

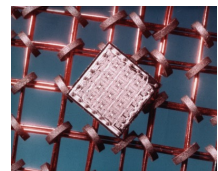
Affinity

HW thrds, cores, sockets



Allocations

API memory routines



Code Transformation

loop transform directives

$for (i, j, k) \rightarrow for(l, m, n)$

OpenMP -- board and membership

The OpenMP Architecture Review Board (ARB)
published first OpenMP Specification in 1997.
It is a non-profit organization with the following members:





OpenMP -- Information

Specifications, Examples and Reference Guide (Cheat Sheet)
and [stackoverflow link](#) are available at:

<https://openmp.org> -- click on Specifications

<https://stackoverflow.com/questions/tagged/openmp>

- Recommended forum for technical questions is [Stack Overflow](#), using the tag OpenMP
- Alternatively post messages on:
 - Twitter: [@OpenMP_ARB](#)
 - LinkedIn: [OpenMP Users](#)
 - Facebook: [OpenMP Group](#)
 - Email: info@openmp.org

<https://www.openmp.org/resources/openmp-compilers-tools/>



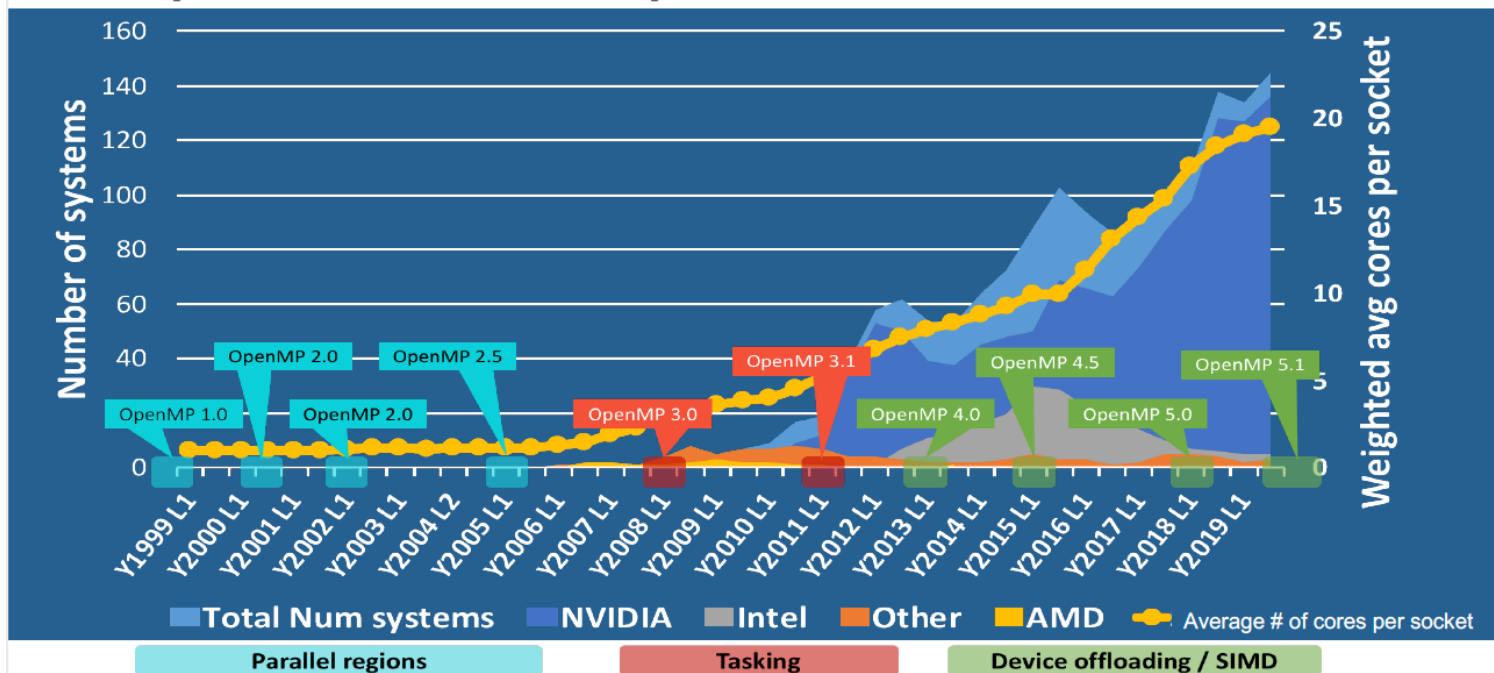
OpenMP Releases

OpenMP Version	Year (Release)	_OPENMP (YYYYMM)	Examples (YY-MM)
5.2	2021	202111	2022-11
5.1	2020	202011	2021-08
5.0	2018	201811	2020-05
4.5	2015	201511	2016-11
4.0	2013	201307	2015-03
3.1	2011	201107	
3.0	2008	200805	



OpenMP -- Evolution

How OpenMP evolves compared with HPC trends (www.top500.org)



Credit: Jose Monsalve Diaz, at University of Delaware

8 Exascale Computing Project



<https://www.openmp.org/resources/openmp-compilers-tools/>



OpenMP -- early development in some areas



2.5→3.0 May 2008

Environment Variables and setter:

Tasking, thread limit, active levels, stack size, wait policy,
set_num_threads, get_ancestor_thread_num, get_levels

3.0→3.1 Jul 2011

Tasking (final, mergeable, taskyield)

Affinity (PROC_BIND)

3.1→4.0 Jul 2013

target directives (data, update, declare, teams, distribute...,)

SIMD directive

User Defined Reduction

Affinity (OMP_PLACES)

Tasking (defined scheduling points, taskgroup, depend clause)

OMP_DISPLAY_ENV

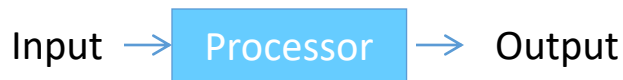


Summary of Basics



Why Parallel?

$$\text{Power} = CV^2F$$

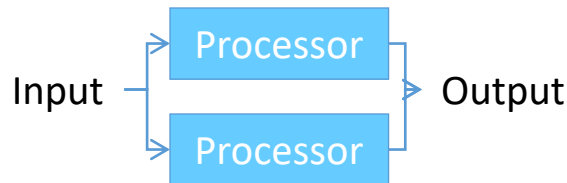


Cap = 1.0c
Volts = 1.0v
Freq = 1.0f
Power = 1.0cv²f

Do same work with 2
processors at 0.5 freq.



Voltage ~ Frequency
2x more wires ->
~2x Capacitance

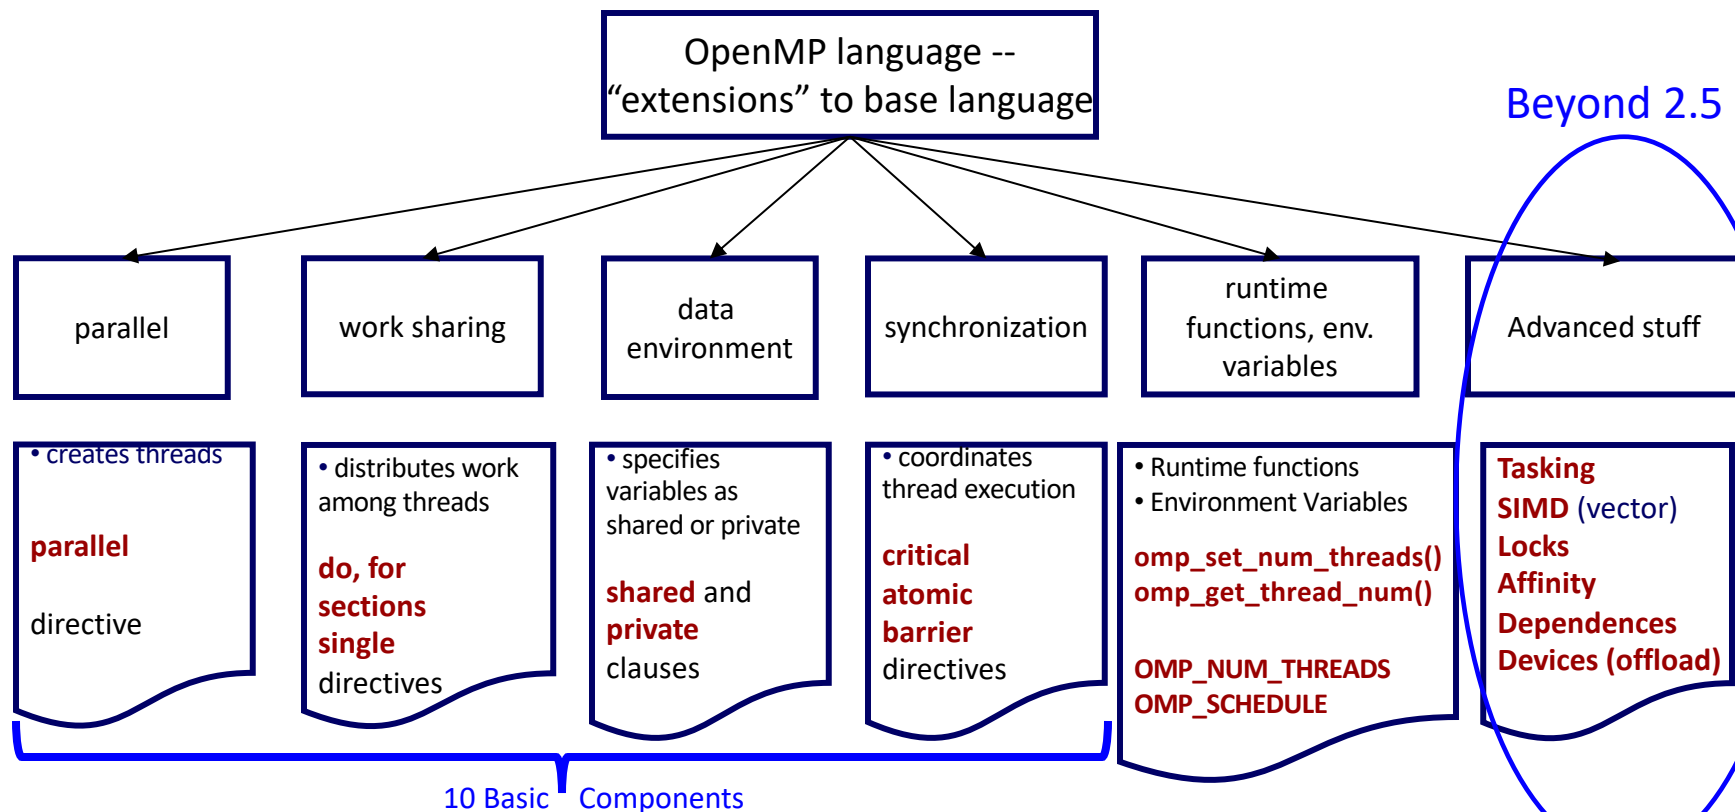


Cap = 2.2c
Volts = 0.6v
Freq = 0.5f
Power = 0.4cv²f

<https://www.youtube.com/watch?v=cMWGeJyrc9w&index=2&list=PL LX-Q6B8xqZ8n8bwjGdzBJ25X2utwn>



OpenMP 2.5 Basics Constructs





OpenMP Syntax

- Compiler directive syntax:

#pragma omp <i>construct</i> [<i>clause</i> [[<i>,</i>] <i>clause</i>]...]	C
!\$omp <i>construct</i> [<i>clause</i> [[<i>,</i>] <i>clause</i>]...]	F90

- Example

Fortran

```
print*, "serial"

!$omp parallel num_threads(4)
...
!$omp end parallel

print*, "serial"
```

C/C++

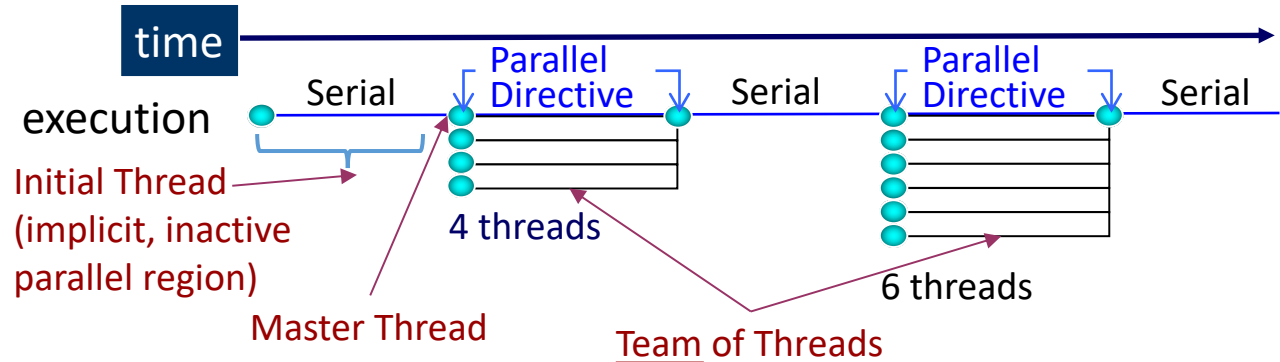
```
printf("serial\n");

#pragma omp parallel num_threads(4)
{
    ... // can be statements or just function call w.o. {}
}

printf("serial\n");
```


Early (2.5 legacy) Execution Model

- Programs begin as a single, *initial thread*.
- A thread encountering a parallel construct it becomes a *master thread*.
- After executing the statements in the parallel region, team threads synchronize and terminate (join) but *initial thread* continues



Note:
[threads](#) – concept of task comes in 3.0.



OpenMP Directive Scope

- OpenMP directives are comments/pragmas in source:

F90 : **!\$omp** free-format*

C/C++ : **#pragma omp** sentinel

* Fortran Fixed Format:
!\$OMP, **C\$OMP** or ***\$OMP**

- Parallel regions are marked by enclosing parallel directives

F90 : **!\$omp parallel ... !\$omp end parallel**

C/C++ : **#pragma omp on block { ... }, or single statement**

- Work-sharing loops are marked by parallel do/for

Fortran

```
!$omp parallel
...
!$omp end parallel

!$omp parallel
  call foo(...)
!$omp end parallel
```

C/C++

```
#pragma omp parallel
{ ...
}

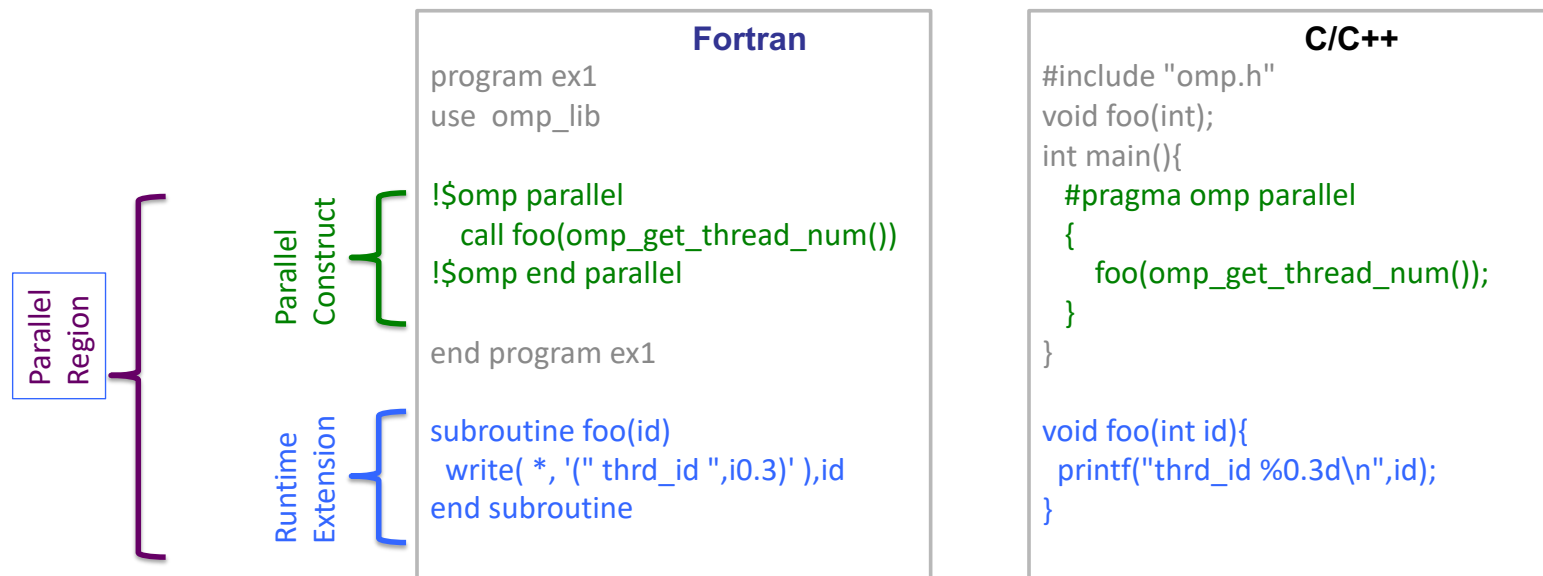
#pragma omp parallel
  foo(...);
```

Any directive that uses **omx** or **ompx**
in the sentinel is implementation defined



OpenMP Directive Scope

- **construct** – the **lexical extent** of executable directive
- **region** – all code encountered in a construct (**construct** + **routines**)





Parallel Region & work ID

F90

For example, to create a 4-thread parallel region.

```
$ export OMP_NUM_THREADS=4; a.out
```

```
real :: A(1000); integer :: ID
!$omp parallel

  ID = omp_get_thread_num()
  call foo(ID, A)
!$omp end parallel
```

But we need to make id
Private to the thread– more later

Each thread executes the code within the structured block

Thread numbers range from 0 to Nthreads-1

Each thread calls foo(ID,A) with a different ID {= 0 to 3}



Parallel Region & work ID

C/C++

For example, to create a 4-thread Parallel region.

```
$ export OMP_NUM_THREADS=4; a.out
```

```
double A[1000];  
#pragma omp parallel  
{  
    int ID = omp_get_thread_num();  
    foo(ID, A);  
}
```

In C++ a local variables are private to the thread...

Each thread executes the code within the structured block

Thread numbers range from 0 to Nthreads-1

Each thread calls foo(ID,A) with a different ID {= 0 to 3}



Parallel Region & Worksharing

Use OpenMP directives to specify Parallel Region & Worksharing constructs

```
...omp parallel  
{  
}  
} $!omp end parallel  
C/C++ F90
```

Code block

Each Thread Executes

do / for
sections
single

Worksharing
Worksharing
Worksharing (one thread)

!\$ and #pragma
not shown here.

Work-sharing Directives **assign threads to units of work**.
There is an **implied barrier at the end of a worksharing** construct!
The *master* directive has no barrier, and is therefore not worksharing.



Parallel Region and Worksharing

```
1  #pragma omp parallel          C/C++
2  {
3      #pragma omp for
4      for (i=0; i<N; i++)
5      {
6          a[i] = b[i] + c[i];
7      }
8  }
```

Line 1 Team of threads formed (parallel region).

Line 3-7 Loop iterations are split among threads.
 implied barrier at }

Each loop iteration must be independent of other iterations.



Parallel Region and Worksharing

```
1  !$OMP PARALLEL                      F90
2      !$OMP DO
3      do i=1,N
4          a(i) = b(i) + c(i)
5      enddo
6  !$OMP END PARALLEL
```

Line 1 Team of threads formed (parallel region).

Line 3-4 Loop iterations are split among threads by DO construct.

Line 5 !\$OMP END DO is optional after the enddo.

Line 5 Implied barrier at enddo.

Each loop iteration must be independent of other iterations.



Parallel Region & Combined Constructs

Fortran

```
!$omp parallel
  !$omp do
    do ...
    end do
  !$omp end do
!$omp end parallel
```

```
!$omp parallel do
  do ...
  end do
```

! *

C/C++

```
# pragma omp parallel
{
  #pragma omp for
  for (...) { ... }
}
```

```
# pragma omp parallel for
for() { ... }
```

* Since Fortran DO has a block structure “!\$omp end do” Is not required

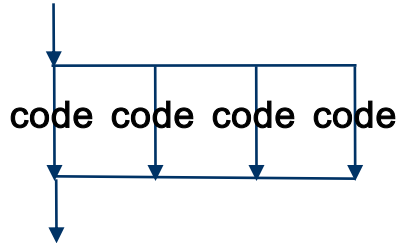
Parallel Region & Combined Constructs

Replicated : Work blocks are executed by all threads.

Work-Sharing : Work is divided among threads.

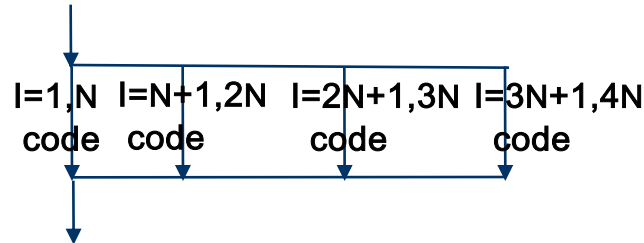
4 Threads

```
!$OMP PARALLEL
{code}
!$OMP END PARALLEL
```



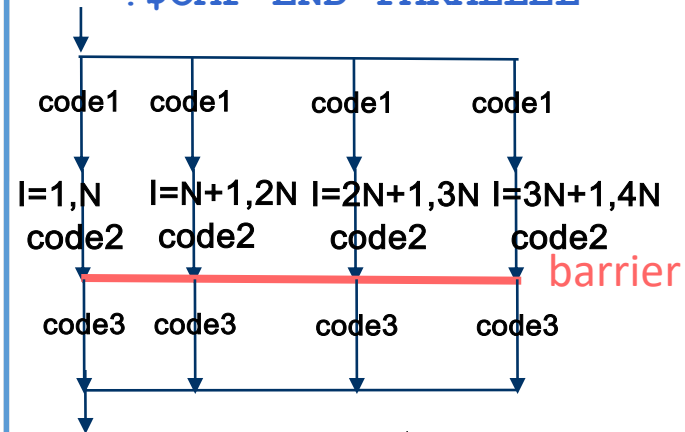
Replicated

```
!$OMP PARALLEL
!$OMP DO
do i = 1,N*4
{code}
end do
!$OMP END PARALLEL DO
```



Work-Sharing

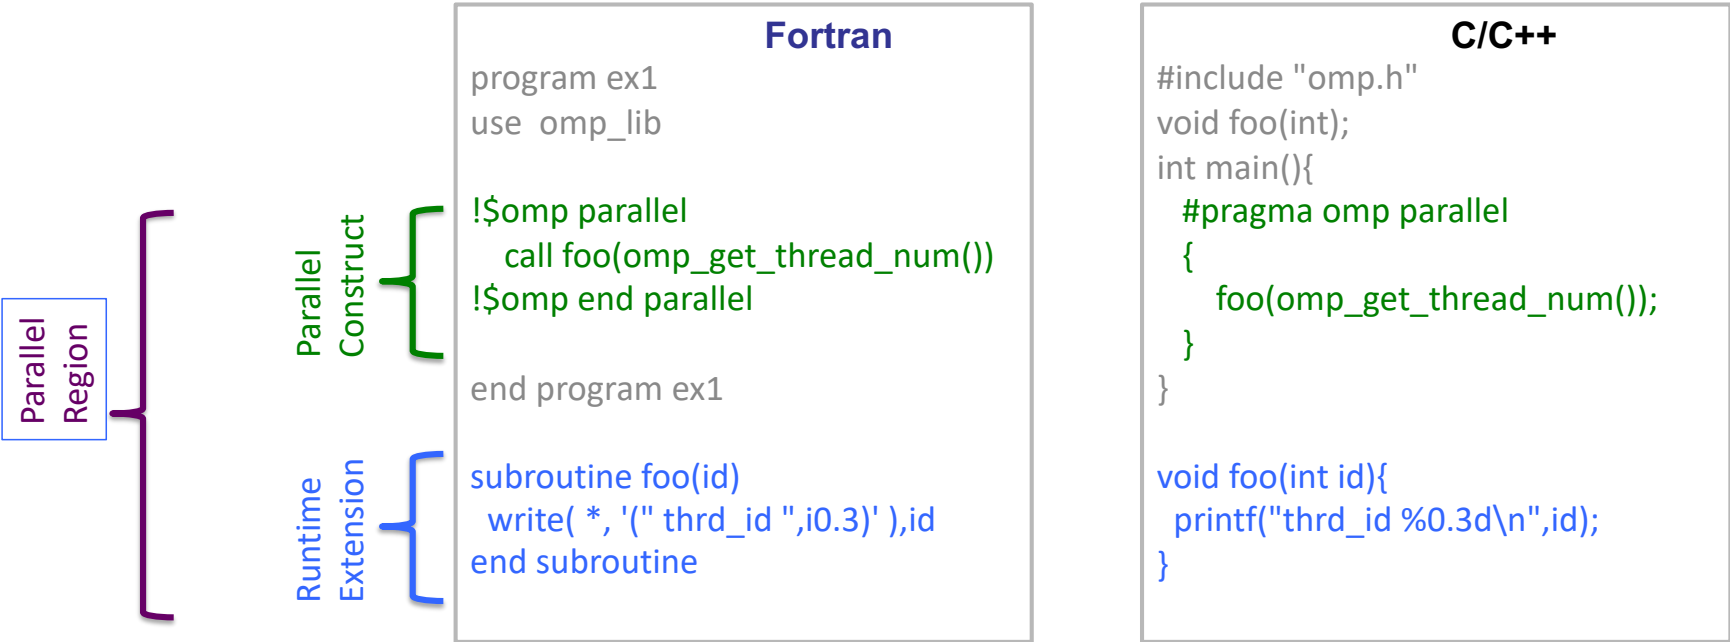
```
!$OMP PARALLEL
{code1}
!$OMP DO
do i = 1,N*4
{code2}
end do
{code3}
!$OMP END PARALLEL
```



Combined

OpenMP Directive Scope

- construct** – the **lexical extent** of executable directive
- region** – all code encountered in a construct (**construct** + **routines**)



Schedule Clause for Worksharing

`schedule(static)`

Each CPU receives one set of contiguous iterations

`schedule(static, C)`

Iterations are divided round-robin fashion in chunks of size C

`schedule(dynamic, C)`

Iterations handed out in chunks of size C as CPUs become available

`schedule(guided, C)`

Each of the iterations are handed out in pieces of exponentially decreasing size, with C minimum number of iterations to dispatch each time

`schedule(runtime)`

Schedule and chunk size taken from the OMP_SCHEDULE environment variable

```
!$omp do schedule(...)
do i=1,128
  A(i)=B(i)+C(i)
enddo
```

```
#pragma omp schedule(
for(i=0;i<128;++i)
  A[i]=B[i]+C[i];
```

`schedule(auto)`

Decision is delegated to compiler and/or runtime

`schedule(monotonic: ...)`

Each thread executes the chunks that it is assigned in increasing logical iteration order (nonmonotonic – any order).

`schedule(simd: ...)` simd size is considered in chunk size—see SIMD.

Kind
Modifier



OpenMP Data Environments

Clauses control the data(-sharing) attributes of vars within a parallel region:

shared, private, reduction, firstprivate, lastprivate

Default variable scope (in parallel region):

1. Variables declared in main/program (C/F90) are shared by default
2. Global variables are shared by default
3. Automatic variables within routines called within a parallel region are private (reside on a stack private to each thread)
4. Loop index of worksharing loops are private.
5. Default scoping rule can be changed with **default** clause



Private and Shared Data

```
!$omp parallel do &  
!$omp& shared(a,b,c,n) private(i)  
  do i = 1,n  
    a(i) = b(i) + c(i)  
  enddo
```

C/C++

```
#pragma omp parallel for \  
  shared(a,b,c,n) private(i)  
  for (i=0; i<n; i++){  
    a[i] = b[i] + c[i];  
  }
```

F90

Reductions

- Operation that combines multiple elements to form a single result
- A variable that accumulates the result is called a reduction variable
- In parallel loops reduction operators and variables must be declared

```
float asum=0.0, aprod=1.0;
#pragma omp parallel for \
    reduction(+:asum) \
    reduction(*:aprod)

for (i=0; i<n; i++){
    asum = asum + a[i];
    aprod = aprod * a[i];
}
```

C/C++

```
real asum=0.0, aprod=1.0
!$omp parallel do          &
!$omp& reduction(+:asum) &
!$omp& reduction(*:aprod)

do i = 1,n
    asum = asum + a(i)
    aprod = aprod * a(i)
enddo
```

F90

Each thread has a private **asum** and **aprod**, initialized to the operator's identity
 After the loop execution, the master thread collects the private values of each thread and finishes the (global) reduction



Synchronization

- Synchronization is used to impose order constraints and to protect access to shared data
- High-Level Synchronization
 - critical
 - atomic **Now has acquire and release semantics!**
 - barrier
 - ordered
- Low-Level Synchronization
 - locks



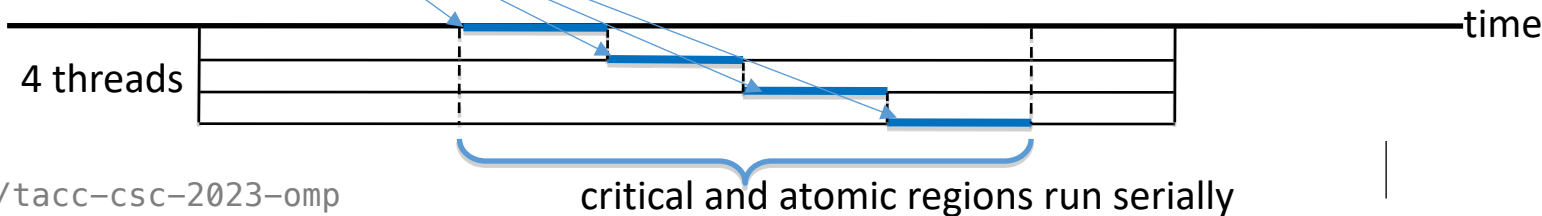
Parallel Region & Combined Constructs

When only one thread at a time can execute a section of code (to avoid race conditions), create a **critical** region. In essence the code executes (and can even run worse than) serially.

Use an **atomic** directive for (“simple”) atomic operations that possibly have hardware atomic support. The expression must have canonical (recognizable) form.

```
#pragma omp parallel shared(sum,x,y)
{ ...
    #pragma omp critical
    {
        update(x);
        update(y);
        sum=sum+1;
    }
}
```

```
!$omp parallel shared(sum,x,y)
...
    !$omp critical
        update(x);
        update(y);
        sum=sum+1;
    !$omp end critical
!$omp end parallel
```



nowait clause (non) synchronization

When a work-sharing region is executed, a barrier is implied - all threads must reach the barrier before any can proceed.

A **nowait** clause at the end of a worksharing region can remove an implied barrier.

```
#pragma omp parallel
{
    #pragma omp for nowait
    {
        for (i=0; i<n; i++)
            {work(i);}
    }
    #pragma omp for schedule(guided,k)
    {
        for (i=0; i<m; i++)
            {x[i]=y[i]+z[i];}
    }
}
```

C/C++

```
!$OMP PARALLEL

    !$OMP DO
        do i=1,n
            work(i)
        enddo
    !$OMP END DO NOWAIT
    !$OMP DO schedule(guided,k)
        do i=1,m
            x(i)=y(i)+z(i)
        enddo
    !$OMP END DO
!$OMP END PARALLEL
```

F90



OpenMP -- Summary of the basics

Concept	What to learn	Level	Optimize
Setup	How to compile OMP_NUM_THREADS	Basic	1 thread per 'core'
Parallel region	Forking/joining threads	Easy	Minimize number of fork/join
Work-sharing/replicated work	What do the threads do? 'omp do/for'	Work-sharing: easy Replicated: medium	Optimize scheduling Remove implicit barriers
Avoiding race conditions	-----	Will take effort!	-----
- Private variables	Why/how to shelter data	medium	
- reduction	Condensing a result from pieces	medium	
- Critical/atomic	All threads, but one thread at a time	harder	Do not serialize everything
- Single/master	One thread, and only one thread	harder	See 'min. fork/join'
Advanced	-----		
- Hybrid	MPI + OpenMP	medium	Interplay MPI/OpenMP
- Thread/memory pinning		medium	Utilize all cores
- SIMD	Vectorization with OpenMP	hard	Utilize vector lanes
- Tasking	Irregular problems	hard	



Hints, Tricks and Tips for OpenMP development



OpenMP Version (Compliance)

Timers

Warm up

OMP_SCHEDULE

Programming style

omp_get... _num_thread(), _thread_num(), etc.



The Compliance Table

Specification and Technical Report	Compliance Version	_OPENMP (YYYYMM)
TR: Nov 2022	for 6.0	
Spec: Nov 2021	5.2	202111
Spec: Nov 2020	5.1	202011
Spec: Nov 2018	5.0	201811
Spec: Nov 2015	4.5	201511
Spec: Jul 2013	4.0	201307
Spec: Jul 2011	3.1	201107

Releases now occur at SuperComputing (in Nov).

Some compilers specify `_OPENMP` with a TR report
--if they have early-adapter implementations

There are many earlier TRs.

If you don't memorize the table, see
openmp.org/specifications where the Spec and TR
links are specified with versions and dates.

See Cray man page `intro_openmp`
(it has details about feature implementation)!



timers -- gettimeofday

Usually accurate to ns.

Call syntax may vary.

Fortran can easily call the C version

```
double gtod_secbase = 0.0E0;
double gtod_timer_() {
    struct timeval tv;
    double sec;

    gettimeofday(&tv, NULL);

    // Always remove the LARGE sec value for improved accuracy
    if(gtod_secbase == 0.0E0)
        gtod_secbase = (double)tv.tv_sec;
    sec = (double)tv.tv_sec - gtod_secbase;

    return sec + 1.0E-06*(double)tv.tv_usec;
}
```



timers -- time stamp counter (tsc)

Highly specific to HW (x86 below, with 2 32-bit registers)

Every Clock Period (CP) counter is updated.

Accuracy is about 20 CP!

But, to get seconds, use the “clock rate”, most often just report CPs.

```
static __inline unsigned long long tsc(void){
    unsigned long a, d;
    unsigned long long d2;

    __asm__ __volatile__ ("rdtsc" : "=a" (a), "=d" (d));

    //                read time stamp counter
    d2 = d;
    return (unsigned long long) a | (d2 << 32);
};

unsigned long long    tsc_timer(void){ return    tsc(); }
unsigned long long    tsc_timer_(void){ return    tsc(); }
```


timers -- class timer

Collects time with a label – good for measuring multiple events.
Easy to use and incorporate.
Uses gettimeofday

```
#include "timer.hpp"
Timer timer;

timer.start("for loop: 1");
    for(...)...;
timer.stop()
    ...
timer.start("function: 2");
    fun();
timer.stop()

timer.print();
```

```
use mod_timer
type(cls_timer) :: timer

call timer%start("do loop: 1")
    do i=1,N ...
call timer%stop
    ...
call timer%start("function: 2")
call foo2()
call timer%stop

call timer%print
```



OpenMP Version (Compliance)

Timers

Warm up

OMP_SCHEDULE

Programming style

omp_get... _num_thread(), _thread_num(), etc.

warm up

- It may take considerable time to get threads from the OS for the first parallel region
- When benchmarking, always request the threads first in a do-nothing parallel region, for doing timings.

```
#include <omp.h>
...
#pragma omp parallel
if(omp_get_num_threads() <
    omp_get_thread_num()) exit(1);

timer.start("parallel loop: 1");
#pragma omp parallel for
for(...)...;
timer.stop();

timer.print();
```

```
use omp_lib
...
!$omp parallel
  if( omp_get_num_threads() < &
      omp_get_thread_num() ) stop

  call timer%start("do loop: 1")
  !$omp parallel do
    do i=1,N ...
  call timer%stop

  call timer%print
```



OMP_SCHEDULE env. var.

- When evaluating performance, you can quickly change the schedule kind by setting the schedule clause to runtime and setting the `OMP_SCHEDULE` environment variable to the kind.

```
$ env OMP_SCHEDULE='static' ./a.out  
$ env OMP_SCHEDULE='dynamic,10' ./a.out
```

```
#pragma omp parallel for \  
    schedule(runtime)  
for(int i=0; i<N; i++) {...}
```

```
$ env OMP_SCHEDULE='static' ./a.out  
$ env OMP_SCHEDULE='dynamic,10' ./a.out
```

```
!$omp parallel do &  
!omp& schedule(runtime)  
do i=N; ...; end do
```



OpenMP Version (Compliance)

Timers

Warm up

OMP_SCHEDULE

Programming style

omp_get... _num_thread(), _thread_num(), etc.

Program styles

- Consistent style can help organize across program files/projects
- Remember, it will be read again by a person and the compiler
 - Use a style that allows a reader to easily evaluate what you have done (and yourself 6 months later)
- Personal experience (KFM) for Fortran coders, for any non-trivial procedure use *implicit none*. (Also, in subroutines use *omp_lib, only : omp_get...* lets readers know what is used.)
- Sometimes, going outside the style box helps in reading and debugging code.

Which "style" is easier to read, understand the logic, and modify?

```
subroutine func2(A)
  implicit none
  double precision  :: A(100,100)
  integer :: i1,i2,j1,j2

  do i1 = 1, 100, 4
    do j1 = 1, 100, 16
      do i2 = i1, i1 + 3
        do j2 = j1, min(j1+15,100)
          A(j2,i2) = A(j2,i2) + 1
        end do
      end do
    end do
  end do
end subroutine
```

```
subroutine func2(A)
  implicit none
  double precision  :: A(100,100)
  integer :: i1,i2,j1,j2

  do i1 = 1, 100, 4
    do j1 = 1, 100, 16
      do i2 = i1, i1 + 3
        do j2 = j1, min(j1+15,100)
          A(j2,i2) = A(j2,i2) + 1
        end do
      end do
    end do
  end do
end subroutine
```



OpenMP API routines

There are many API routines that can help you

- * Threads and Teams: Count, number(ID), max -- getters and setters
- * Target: Device count, number, Is Initial Device
- * Parallel: In parallel region
- * Memory allocation routines (host and device)
- * Affinity Information – display and setters

Most are self descriptive:



OpenMP APIs

Most are self
descriptive:

omp_set_num_threads
omp_get_num_threads
omp_get_max_threads
omp_get_thread_num
omp_in_parallel

omp_set_schedule
omp_get_schedule
omp_get_thread_limit
omp_get_team_size

omp_get_proc_bind
omp_get_num_places
omp_get_place_num_procs

...
omp_set_affinity_format
omp_get_affinity_format
omp_display_affinity
omp_capture_affinity

omp_get_num_teams
omp_get_team_num
omp_set_num_teams
omp_get_max_teams
omp_set_teams_thread_limit
omp_get_teams_thread_limit
omp_get_num_procs
omp_set_default_device
omp_get_default_device
omp_get_num_devices
omp_get_device_num
omp_is_initial_device
omp_get_initial_device

omp_target_alloc
omp_target_free
omp_target_is_present
omp_target_is_accessible
omp_target_memcpy
omp_target_memcpy_rect
omp_target_memcpy_async
omp_target_memcpy_rect_async
omp_target_associate_ptr
omp_target_disassociate_ptr
omp_get_mapped_ptr

omp_init_lock
omp_init_lock_with_hint
omp_init_nest_lock_with_hint
omp_destroy_lock
omp_set_lock
omp_unset_lock
omp_test_lock

omp_get_wtime
omp_get_wtick

omp_init_allocator
omp_destroy_allocator
omp_set_default_allocator
omp_get_default_allocator
omp_alloc
omp_free
omp_calloc
omp_realloc

.....NOT all included!



- DISPLAY Environment Variable (EV) info

```
export OMP_DISPLAY_ENV=TRUE | FALSE
```

Shows ICVs associated with the EVs
after runtime evaluates EVs (or their defaults).

Use **VERBOSE** to show values of runtime
variables that may be modified by vendor-specific.

*ICV = internal control variables associated with a data environment

- 5.1

API routine:

```
omp_display_env(verbose) //C/C++  
call omp_display_env(verbose) !! F90
```



Program Information

OPENMP DISPLAY ENVIRONMENT BEGIN

`_OPENMP='202011'`

[host] `OMP_NUM_THREADS`: value is not defined

[host] `OMP_MAX_ACTIVE_LEVELS`='1'

[host] `OMP_SCHEDULE`='static'

[host] `OMP_STACKSIZE`='4M'

[host] `OMP_DISPLAY_AFFINITY`='FALSE'

[host] `OMP_AFFINITY_FORMAT`='(null)'

[host] `OMP_PLACES`: value is not defined

[host] `OMP_PROC_BIND`: value is not defined

[host] `OMP_TARGET_OFFLOAD`=DEFAULT

[host] `OMP_DEFAULT_DEVICE`='0'

[host] `OMP_ALLOCATOR`='omp_default_mem_alloc'

[host] `OMP_MAX_TASK_PRIORITY`='0'

OPENMP DISPLAY ENVIRONMENT END

`_OPENMP: YYYYMM`

Year & Month of Spec.

201511 – 4.5

201811 – 5.0

202011 – 5.1

202111 – 5.2

Some vendors are
Technical Report
compliant

202008 – TR9

202107 – TR10

202211 – TR11

Not all ENV VARs shown.
Order has been changed.



Program Information Displays

- Display Affinity Info

ENV. VAR:

```
export OMP_DISPLAY_AFFINITY=TRUE | FALSE
```

API routine:

```
omp_display_affinity(NULL) //CIC++  
call omp_display_affinity("") !! F90
```

Can put format string here.

Wall Clock Timer

- Every developer should have an easy way to evaluate performance for code optimization and changes
- Total time:
 wrapper: `/usr/bin/time -p ./a.out`
 wrapped: `date +%s`
 `./a.out`
 `date +%s #in secs.`

```
$ /usr/bin/time -p a.out
real 5.00
user 4.50
sys  0.01
```

```
$ date; ./a.out; date
1693573722
1693573727
```

Bash function nanosecond timer:

```
nstimer(){ t0=`date +%s.%N`; $*; t1=`date +%s.%N`; bc <<<"$t1 - $t0" ; }
```

```
$ nstimer echo Hello World
Hello World
.002101725
```

thread number

- API routine: `omp_get_thread_num()`
- In normal programming used for task-parallel work assignment in a parallel region.
- Identifying threads is useful for (debugging):
 - ✓ Making sure multiple threads are executing tasks
 - ✓ Critical in resolving affinity problems
 - ✓ Resolving race conditions by slowing down one thread in a race.



primary thread

An OpenMP thread that has thread number 0. A primary thread may be an initial thread or the thread that encounters a parallel construct, creates a team, generates a set of implicit tasks, and then executes one of those tasks as thread number 0.

Example of Class Timer for C/C++/F90 codes

- See timer.hpp and timer.f90

```
#include "timer.hpp";
int main(){
    Timer timer;

    timer.start("CPU:foo1");
    foo1();
    timer.stop();

    timer.start("CPU:foo2");
    foo2();
    timer.stop();

    timer.print();
}
```

```
#include "timer.f90"
program main
    use mod_timer
    type(cls_timer) :: timer
    call timer%start("CPU:foo1")
    call foo1()
    call timer%stop

    call timer%start("CPU:foo2")
    call foo2()
    call timer%stop

    call timer%print
end
```



Output from Class Timer for C/C++/F90 codes

- OUTPUT:

```
Action      ::      time/s      Time resolution = 1e-06
-----
CPU:foo1    ::      1.9e-05
CPU:foo2    ::      1.0e-06
```


- `|& tee file` for displaying stdout/err and saving to *file*
- *Contributions welcome:* _____

Remember how to get to these references

- OpenMP References

<https://www.openmp.org>

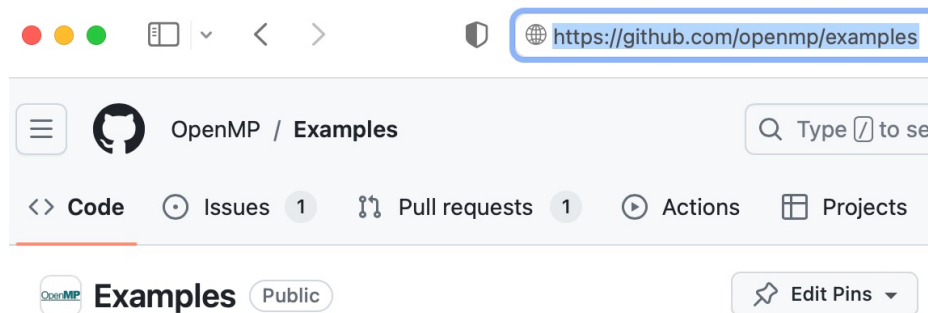
Look under Specifications

[OpenMP API 5.2 Specification](#)

[OpenMP API 5.2.1 Examples](#)

[OpenMP API 5.2 Reference Guide](#)

- <https://github.com/openmp/examples> get example codes here





For Exercises, git clone

<https://github.com/csc-training/advancedOpenMP.git>

```
$ cd OpenMP_cpu
```

```
$ cd adv_intro
```

```
$ ls
```

```
1_env_apis
```

```
2_timers
```

```
3_private_if_clauses
```

```
README_dev_access.md
```

```
README.md
```

```
job_lumi_C
```